

Многонишково програмиране с JAVA

Multithreading

- За какво е необходимо многонишково програмиране?
 - Програманите езици са последователни, т.е. инструкциите се изпълняват една след друга(ред след ред)

```
public static void main(String[] args) {  
  
    initializeArrays();  
    downloadStocks();  
    initializeTimeSeriesModels();  
    makePredictions();  
  
}
```

- При еднонишкова програма тези методи ще бъдат извикани един след друг: трябва да се изчака изпълнението да свърши едно по едно
 - Трябва да се отбележи, че не е най-доброто решение: операциите, които консумират много време трябва да се изчакват и потребителите няма да знаят какво се случва!

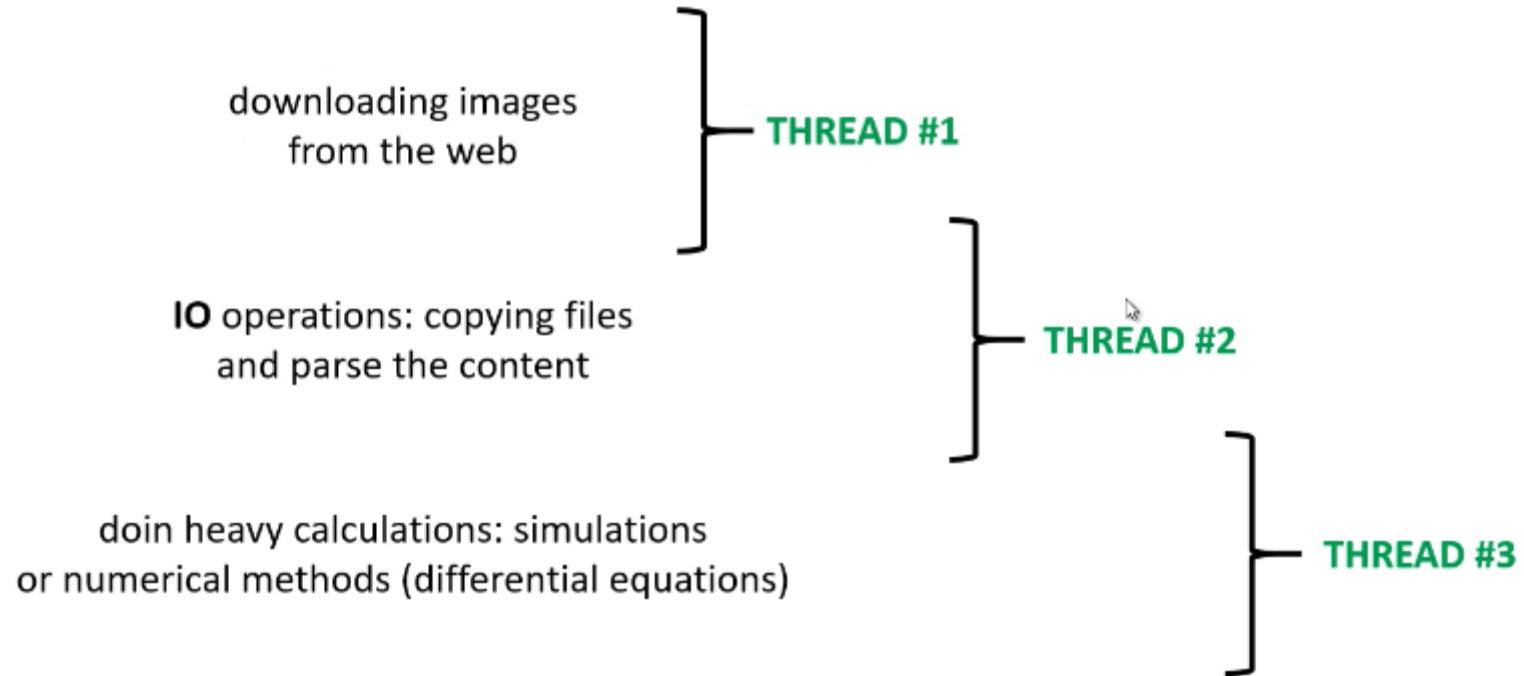
Multithreading

- Основният принцип в multithreading е да се раздели изпълнението на подзадачи: всяка една или група от подзадачи да се изпълняват едновременно.

For example: our stock market related software is able to download data from the web (*Yahoo Finance* for example)

- it takes **2-3** mins to fetch the data BUT we want to make sure the application is not frozen !!!
- solution: we create a distinct thread for the download operation and during this procedure the user can do whatever he/she wants in the application

Multithreading



Multithreading

„Multithreading is the ability of the **CPU** to execute multiple processes or threads concurrently”

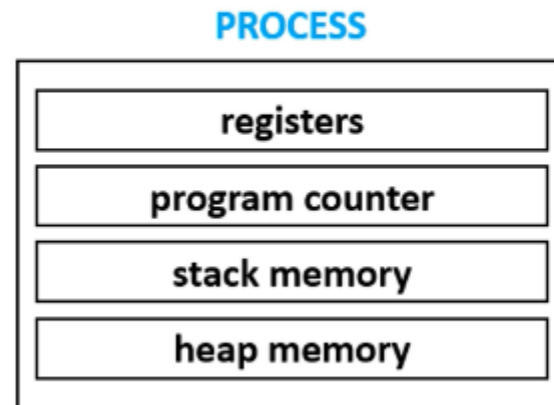
Both **processes** and **threads** are independent sequences of execution

PROCESS: a process is an instance of program execution

→ when you open a software or a web browser: these are distinct processes

→ the **OS** assigns distinct registers, stack memory and heap memory to every single process

in **Java** we can create
processes with the help of
ProcessBuilder class



Microsoft Windows [Version 10.0.18362.476]
(c) 2019 Microsoft Corporation. All rights reserved.

C:\Users\Tianhe-2>tasklist

| Image Name | PID | Session Name | Session# | Mem Usage |
|---------------------|------|--------------|----------|-----------|
| System Idle Process | 0 | Services | 0 | 8 K |
| System | 4 | Services | 0 | 1 696 K |
| Secure System | 72 | Services | 0 | 28 884 K |
| Registry | 128 | Services | 0 | 78 068 K |
| smss.exe | 408 | Services | 0 | 1 096 K |
| csrss.exe | 736 | Services | 0 | 4 196 K |
| wininit.exe | 856 | Services | 0 | 5 372 K |
| csrss.exe | 864 | Console | 1 | 5 044 K |
| services.exe | 928 | Services | 0 | 9 556 K |
| lsass.exe | 948 | Services | 0 | 2 580 K |
| lsass.exe | 956 | Services | 0 | 19 764 K |
| winlogon.exe | 1020 | Console | 1 | 12 892 K |
| svchost.exe | 628 | Services | 0 | 3 544 K |
| svchost.exe | 1044 | Services | 0 | 30 596 K |
| fontdrvhost.exe | 1068 | Console | 1 | 7 052 K |
| fontdrvhost.exe | 1072 | Services | 0 | 2 200 K |
| WUDFHost.exe | 1132 | Services | 0 | 8 196 K |
| svchost.exe | 1208 | Services | 0 | 17 976 K |
| svchost.exe | 1264 | Services | 0 | 9 244 K |
| dwm.exe | 1348 | Console | 1 | 104 360 K |
| svchost.exe | 1428 | Services | 0 | 13 224 K |

Multithreading

„Multithreading is the ability of the **CPU** to execute multiple processes or threads concurrently”

Both **processes** and **threads** are independent sequences of execution

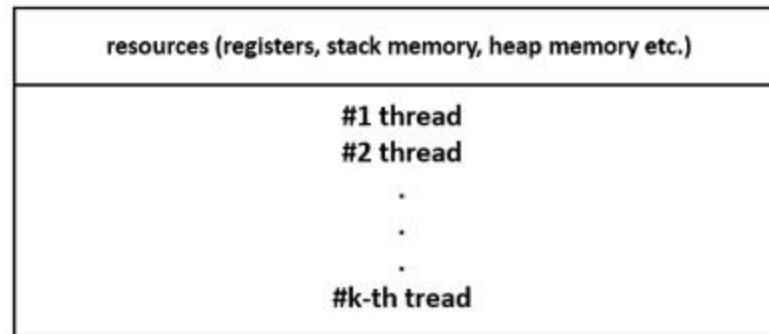
THREADS: a thread is a „light-weight” process

→ It is a unit of execution within a given process (a process may have several threads)

→ each thread in a process shares the memory and resources and this is why programmers have to deal with concurrent programming and multithreading

PROCESS

creating a new thread
requires fewer resources
than creating a new process



Multithreading

MULTITHREADING AND TIME-SLICING ALGORITHM

Assume we have **k** threads (so more than **1** thread in our application)

→ somehow the single processor has to deal with all of the **k** threads
~ one approach is to use time-slicing algorithm

„Processing time for a single core is shared among processes and threads. This is called time-slicing”



„multithreaded execution” (with time-slicing)

In this case the **CPU** will run **thread #1** for a small amount of time then **thread #2** then again **thread #1** and then **thread #2** and so on ...

Multithreading

- **Предимства**

- Създават се по-бързи приложения: може да представят няколко операции едновременно

- Постига се по-добра ресурсна натовареност (**CPU utilization**)

По подразбиране всяка програма на Java е еднонишкова: има няколко ядра, които няма да се използват без многонишков режим.

- Може да подобри производителността

// подобрението на производителността ще има, когато се използват множество ядра и паралелно програмиране


Multithreading

MULTITHREADING AND TIME-SLICING ALGORITHM

Assume we have **k** threads (so more than **1** thread in our application)

→ somehow the single processor has to deal with all of the **k** threads
~ one approach is to use time-slicing algorithm


„Processing time for a single core is shared among processes and threads. This is called time-slicing“



In this case the **CPU** will run **thread #1** for a small amount of time then **thread #2** then again **thread #1** and then **thread #2** and so on ...

Multithreading

DISADVANTAGES



It's expensive to switch between threads: this is why using threads is not always the fastest way possible (for example sorting algorithms)

RULE OF THUMB: for small problems it is unnecessary to use multithreading

Multithreading

THREAD STATES

- 1.) **NEW** when we instantiate a thread
It is in this state until we start it
~ **start()** method
- 2.) **RUNNABLE** after we have started the thread
The thread is executing its task in this state
- 3.) **WAITING** when a thread is waiting: for example for another thread to finish its task
When other thread signals then this thread goes back to the **runnable** state
~ **wait()** and **sleep()** methods
- 4.) **DEAD** after the thread finishes its task

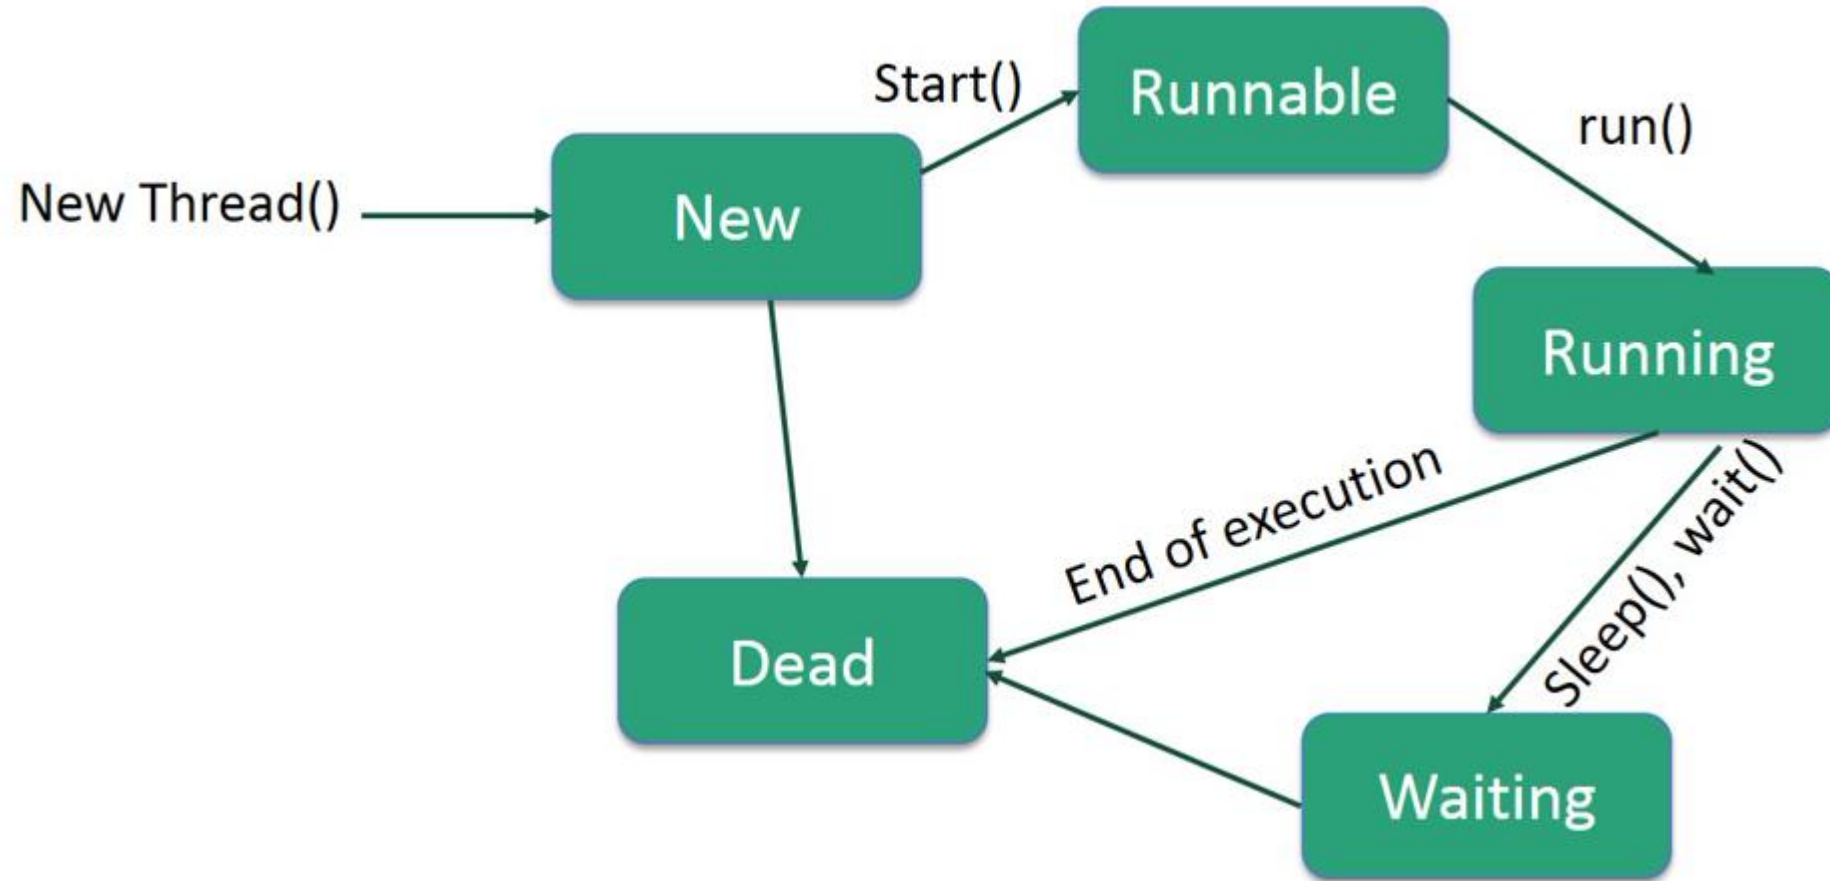
```
1 package com.balazsholczer.udemy;
2
3 public class App {
4
5     private static int counter = 0;
6
7     public static synchronized void increment() {
8         ++counter;
9     }
10
11     public static void process() {
12
13         Thread t1 = new Thread(new Runnable() {
14
15             @Override
16             public void run() {
17                 for (int i = 0; i < 100; ++i)
18                     increment();
19             }
20         });
21
22         Thread t2 = new Thread(new Runnable() {
23
24             @Override
25             public void run() {
26                 for (int i = 0; i < 100; ++i)
27                     increment();
28             }
29         });
30
31         t1.start();
```

```
10
11 public static void process() {
12
13     Thread t1 = new Thread(new Runnable() {
14
15         @Override
16         public void run() {
17             for (int i = 0; i < 100; ++i)
18                 increment();
19         }
20     });
21
22     Thread t2 = new Thread(new Runnable() {
23
24         @Override
25         public void run() {
26             for (int i = 0; i < 100; ++i)
27                 increment();
28         }
29     });
30
31     t1.start();
32     t2.start();
33
34     try {
35         t1.join();
36         t2.join();
37     } catch (InterruptedException e) {
```

Multithreading

THREAD STATES

- 1.) **NEW** when we instantiate a thread
It is in this state until we start it
~ **start()** method
- 2.) **RUNNABLE** after we have started the thread
The thread is executing its task in this state
- 3.) **WAITING** when a thread is waiting: for example for another thread to finish its task
When other thread signals then this thread goes back to the **runnable** state
~ **wait()** and **sleep()** methods
- 4.) **DEAD** after the thread finishes its task




```
7 public static synchronized void increment() {
8     ++counter;
9 }
10
11 public static void process() {
12
13     Thread t1 = new Thread(new Runnable() {
14
15         @Override
16         public void run() {
17             for (int i = 0; i < 100; ++i)
18                 increment();
19         }
20     });
21
22     Thread t2 = new Thread(new Runnable() {
23
24         @Override
25         public void run() {
26             for (int i = 0; i < 100; ++i)
27                 increment();
28         }
29     });
30
31     t1.start();
32     t2.start();
33
34     try {
35         t1.join();
36         t2.join();
37     } catch (InterruptedException e) {
```

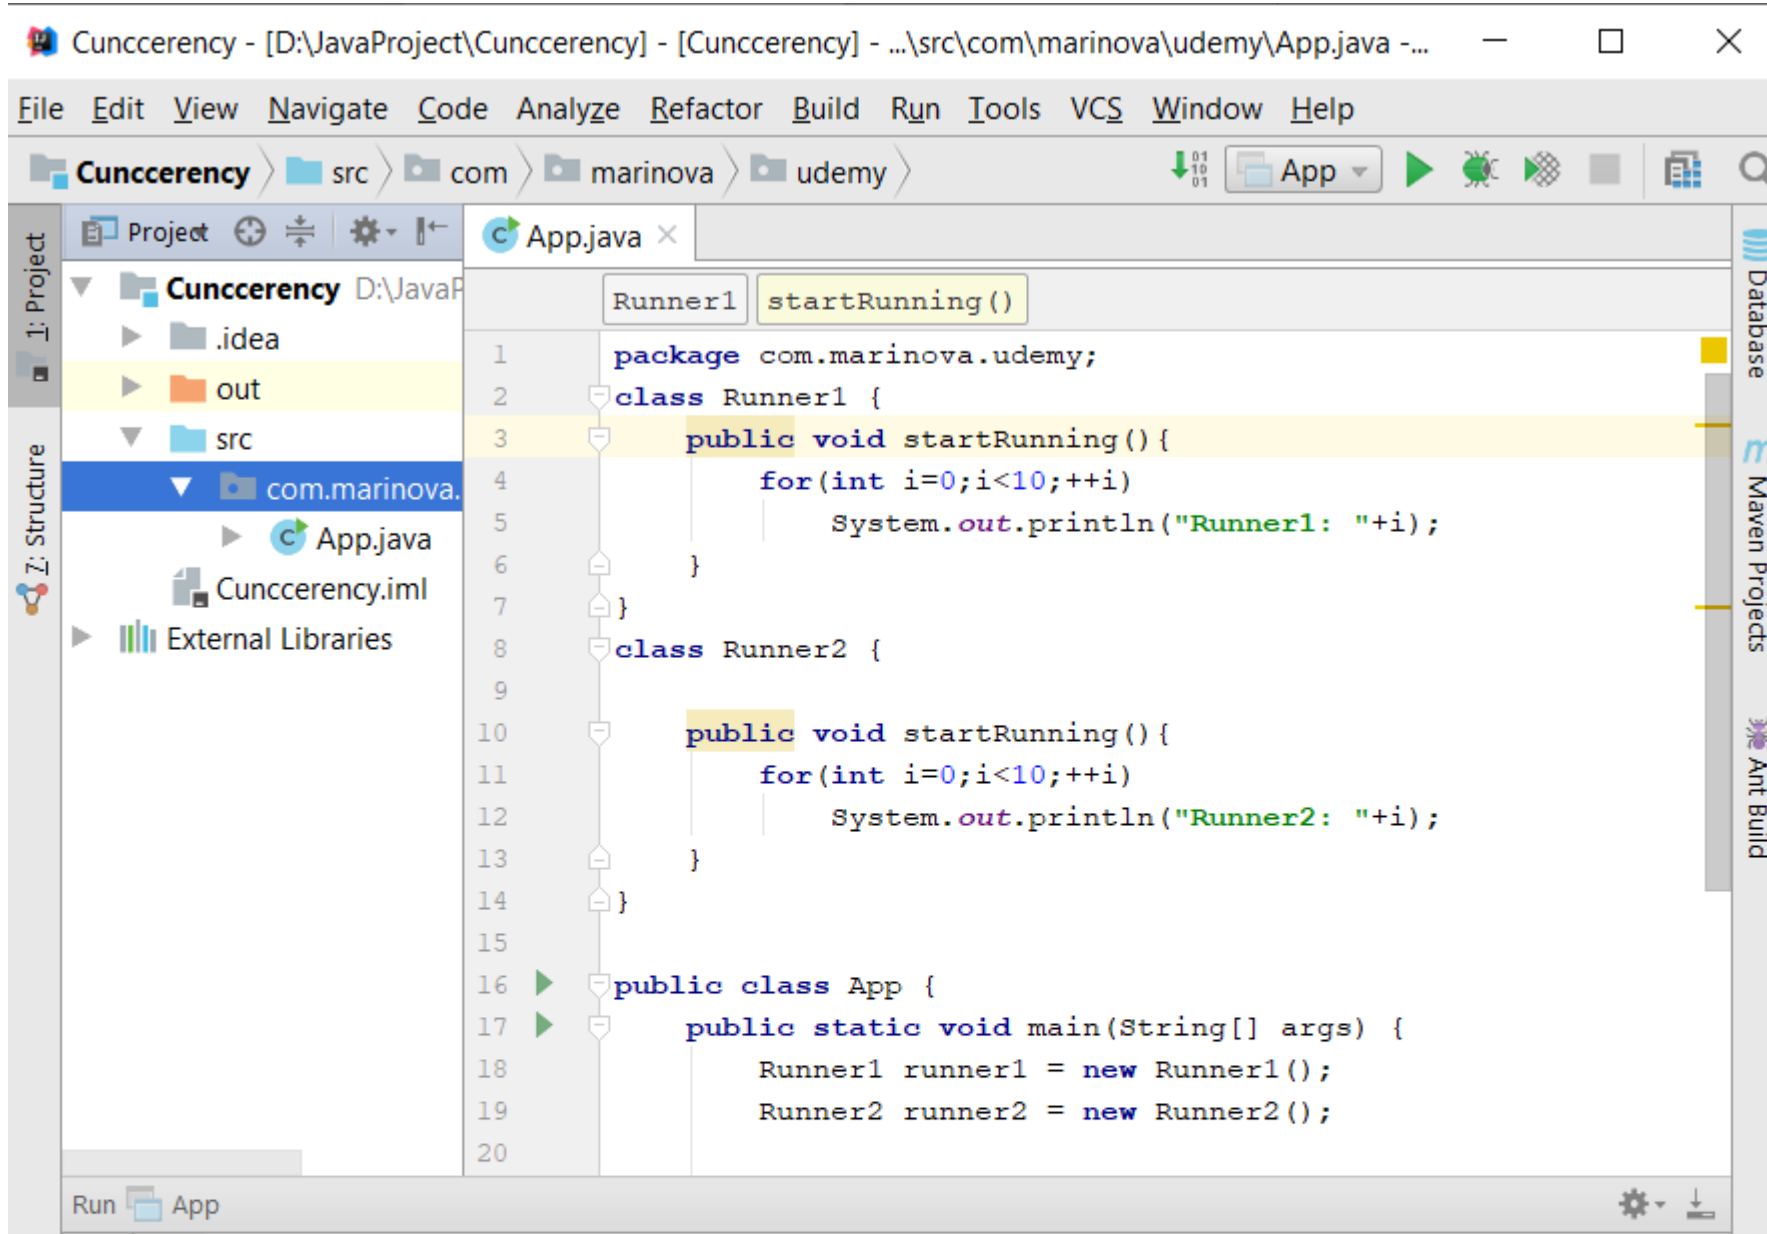
```
7 public static synchronized void increment() {
8     ++counter;
9 }
10
11 public static void process() {
12
13     Thread t1 = new Thread(new Runnable() {
14
15         @Override
16         public void run() {
17             for (int i = 0; i < 100; ++i)
18                 increment();
19         }
20     });
21
22     Thread t2 = new Thread(new Runnable() {
23
24         @Override
25         public void run() {
26             for (int i = 0; i < 100; ++i)
27                 increment();
28         }
29     });
30
31     t1.start();
32     t2.start();
33
34     try {
35         t1.join();
36         t2.join();
37     } catch (InterruptedException e) {
```

Multithreading

THREAD STATES

- 1.) **NEW** when we instantiate a thread
It is in this state until we start it
~ **start()** method
- 2.) **RUNNABLE** after we have started the thread
The thread is executing its task in this state
- 3.) **WAITING** when a thread is waiting: for example for another thread to finish its task
When other thread signals then this thread goes back to the **runnable** state
~ **wait()** and **sleep()** methods
- 4.) **DEAD** after the thread finishes its task

Sequential processing



Резултат от изпълнението на последователната програма:

```
"C:\Program Files\Java\jdk1.8.0_221\bin\java" ...
```

```
Runner1: 0
```

```
Runner1: 1
```

```
Runner1: 2
```

```
Runner1: 3
```

```
Runner1: 4
```

```
Runner1: 5
```

```
Runner1: 6
```

```
Runner1: 7
```

```
Runner1: 8
```

```
Runner1: 9
```

```
Runner2: 0
```

```
Runner2: 1
```

```
Runner2: 2
```

```
Runner2: 3
```

```
Runner2: 4
```

```
Runner2: 5
```

```
Runner2: 6
```

```
Runner2: 7
```

```
Runner2: 8
```

```
Runner2: 9
```

```
Process finished with exit code 0
```

Project: 1: Project

Structure: Cunccerency D:\JavaP

- App
- main()
- App
- main()

```
1 package com.marinova.udemy;
2 class Runner1 implements Runnable {
3
4     @Override
5     public void run() {
6         for(int i=0;i<10;++i)
7             System.out.println("Runner1: "+i);
8     }
9 }
10 class Runner2 implements Runnable {
11
12     @Override
13     public void run() {
14         for(int i=0;i<10;++i)
15             System.out.println("Runner2: "+i);
16     }
17 }
18
19 public class App {
20     public static void main(String[] args) {
21
22         Thread t1 = new Thread((new Runner1()));
23         Thread t2 = new Thread((new Runner2()));
24
25         t1.start();
26         t2.start();
27     }
28 }
29
```

Database

Maven Projects

Ant Build

Резултат от multithreading изпълнение:



The screenshot shows an IDE's console window with tabs for Problems, Javadoc, Declaration, and Console. The console output displays the execution of a Java application with two threads, Runner1 and Runner2. The output is as follows:

```
<terminated> App [Java Application] C:\Program Files\Java\jre1.8
Runner1: 0
Runner2: 0
Runner1: 1
Runner2: 1
Runner1: 2
Runner2: 2
Runner1: 3
Runner2: 3
Runner1: 4
Runner2: 4
Runner1: 5
Runner2: 5
Runner2: 6
Runner2: 7
Runner2: 8
Runner2: 9
Runner1: 6
Runner1: 7
Runner1: 8
Runner1: 9|
```

The output demonstrates that Runner2 completes its execution (values 0 through 9) before Runner1 (values 0 through 9), illustrating the effects of multithreading.

```

package com.marinova.udemy;
class Runner1 extends Thread {

    @Override
    public void run() {
        for(int i=0;i<10;++i)
            System.out.println("Runner1: "+i);
    }
}

class Runner2 extends Thread {

    @Override
    public void run() {
        for(int i=0;i<10;++i)
            System.out.println("Runner2: "+i);
    }
}

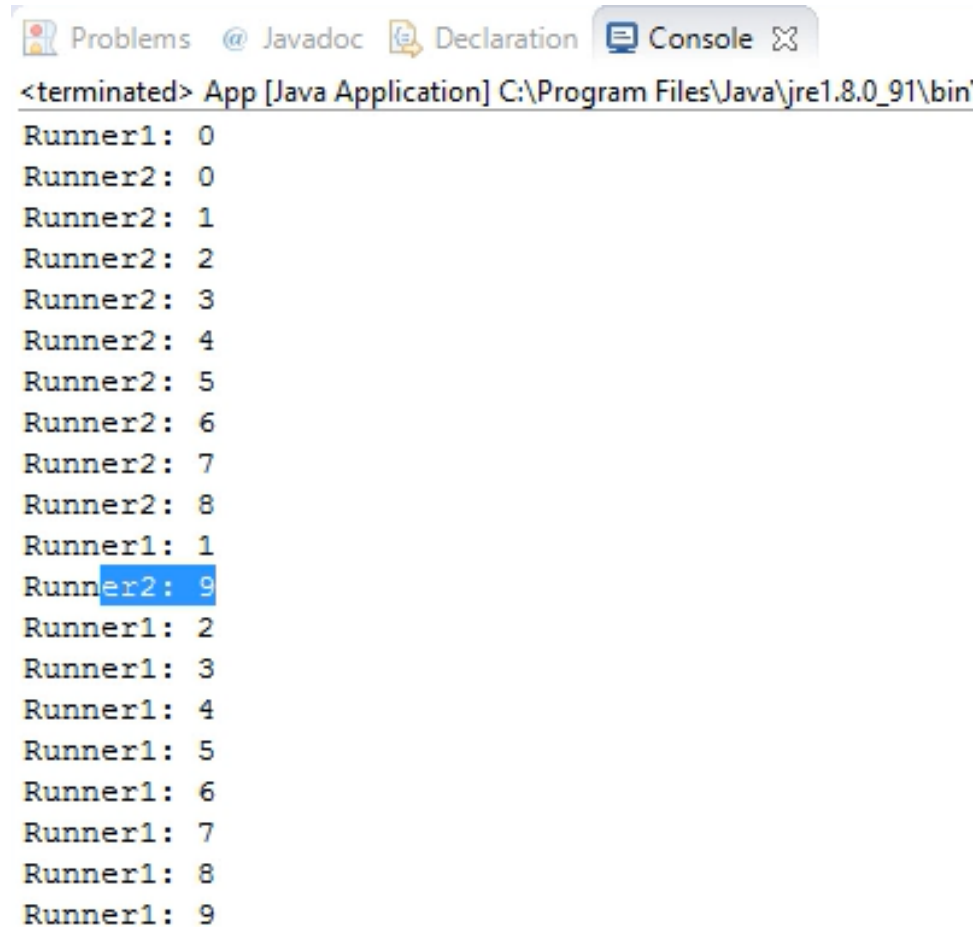
public class App {
    public static void main(String[] args) {

        /*      Thread t1 = new Thread((new Runner1()));
               Thread t2 = new Thread((new Runner2()));*/
        Runner1 t1 = new Runner1();
        Runner2 t2 = new Runner2();
        t1.start();
        t2.start();
    }
}

```

- Друг вариант – наследяване на класа Thread


Резултат от изпълнението:



```
<terminated> App [Java Application] C:\Program Files\Java\jre1.8.0_91\bin'
Runner1: 0
Runner2: 0
Runner2: 1
Runner2: 2
Runner2: 3
Runner2: 4
Runner2: 5
Runner2: 6
Runner2: 7
Runner2: 8
Runner1: 1
Runner2: 9
Runner1: 2
Runner1: 3
Runner1: 4
Runner1: 5
Runner1: 6
Runner1: 7
Runner1: 8
Runner1: 9
```

- Допълнителна задача: Стартирайте програмата, така че броя на итерациите да е 100. Разгледайте каква е последователността на изпълнение на t1 и t2.

```
package com.marinova.udemy;
class Runner1 extends Thread {

    @Override
    public void run() {
         for(int i=0;i<100;++i){
            System.out.println("Runner1: "+i);
            try {
                Thread.sleep( millis: 100);
            } catch (InterruptedException e) {
                e.printStackTrace();
            }
        }
    }
}

class Runner2 extends Thread {

    @Override
    public void run() {
        for (int i = 0; i < 100; ++i) {
            System.out.println("Runner2: " + i);
            try {
                Thread.sleep( millis: 100);
            } catch (InterruptedException e) {
                e.printStackTrace();
            }
        }
    }
}
```

```
"C:\Program Files\Java\jdk1.8.0_221\bin\java" ...
```

Runner2: 0

Runner1: 0

Runner1: 1

Runner2: 1

Runner2: 2

Runner1: 2

Runner2: 3

Runner1: 3

Runner1: 4

Runner2: 4

Runner2: 5

Runner1: 5

Runner1: 6

Runner2: 6

Runner1: 7

Runner2: 7

Runner1: 8

Runner2: 8

Runner1: 9

Runner2: 9

Runner2: 10

Runner1: 10

Runner1: 11

Runner2: 11

Runner2: 12

Runner1: 12

Runner2: 13

Runner1: 13

Runner1: 14


Runner2: 14

Runner2: 15

Runner1: 15

Runner2: 16

Parallel versus multithreading



Runnable vs Thread

1. Create Thread using Runnable Interface vs Thread class

1.1. Runnable interface

```
public class DemoRunnable implements Runnable {  
    public void run() {  
        //Code  
    }  
}  
  
//start new thread with a "new Thread(new demoRunnable()).start()" call
```

Runnable vs Thread – cont.

1.2. Thread class

```
public class DemoThread extends Thread {  
    public DemoThread() {  
        super("DemoThread");  
    }  
    public void run() {  
        //Code  
    }  
}  
//start new thread with a "new demoThread().start()" call
```

Runnable vs Thread – cont.


2. Difference between Runnable vs Thread

```
33 public class App {
34
35     public static void main(String[] args) {
36
37         // Thread t1 = new Thread(new Runner1());
38         // Thread t2 = new Thread(new Runner2());
39
40         Runner1 t1 = new Runner1();
41         Runner2 t2 = new Runner2();
42
43         t1.start();
44         t2.start();
45
46         System.out.println("Finished the tasks...");
47     }
48 }
49
```

Problems @ Javadoc Declaration Console

```
<terminated> App [Java Application] C:\Program Files\Java\jre1.8.0_91\bin\javaw.exe (2)
Finished the tasks...
Runner2: 0
Runner1: 0
Runner2: 1
Runner1: 1
Runner2: 2
```


Volatile



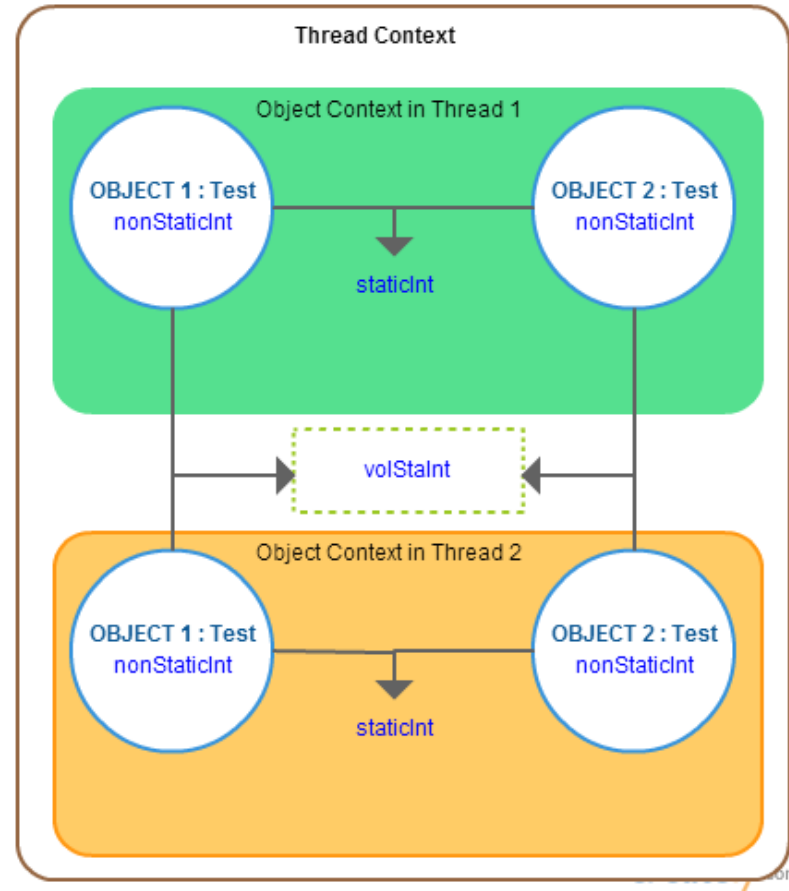
Every read of a volatile variable will be read from the **RAM** so from the main memory

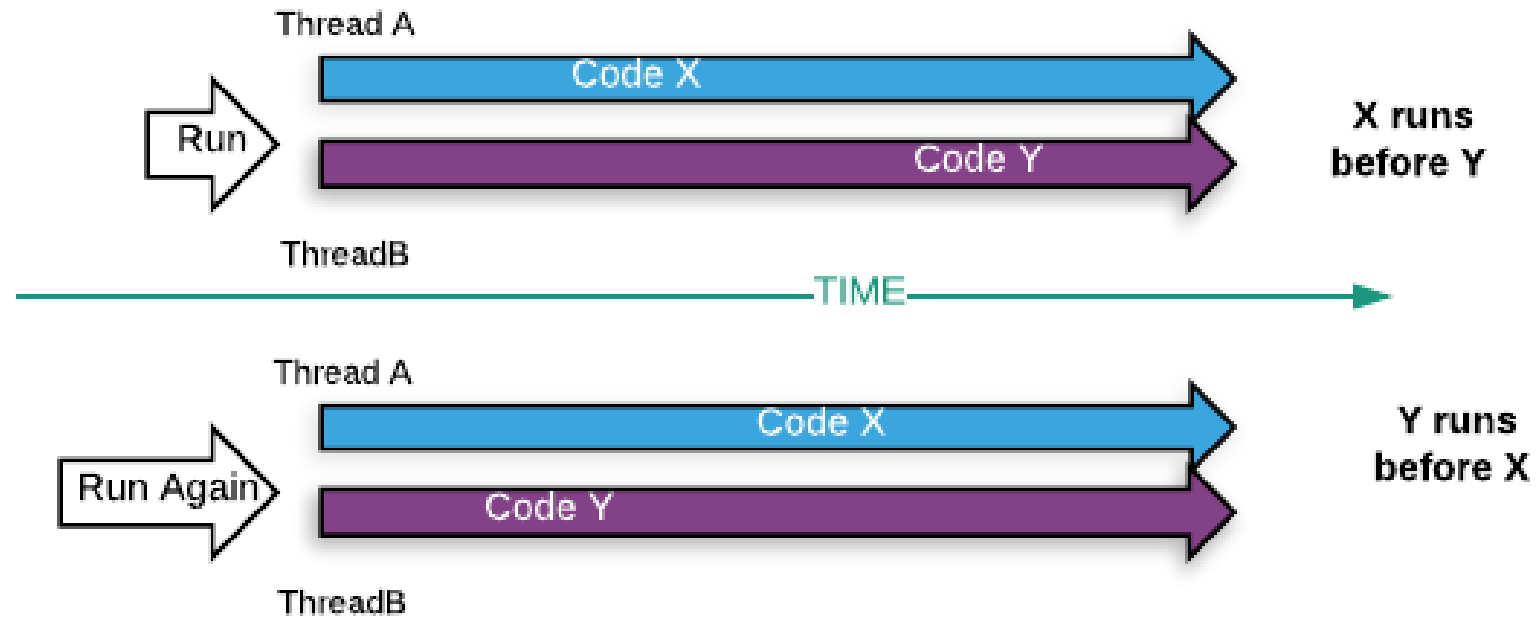
// not from cache, usually variables are cached for performance reasons

Caches are faster → do not use **volatile** keyword if not necessary
(+ it prevents instruction reordering which is a performance boost technique)

Volatile Vs Static in JAVA

```
Class Test {  
    int nonStaticInt;  
    static int staticInt;  
    volatile static int volStaticInt;  
}
```

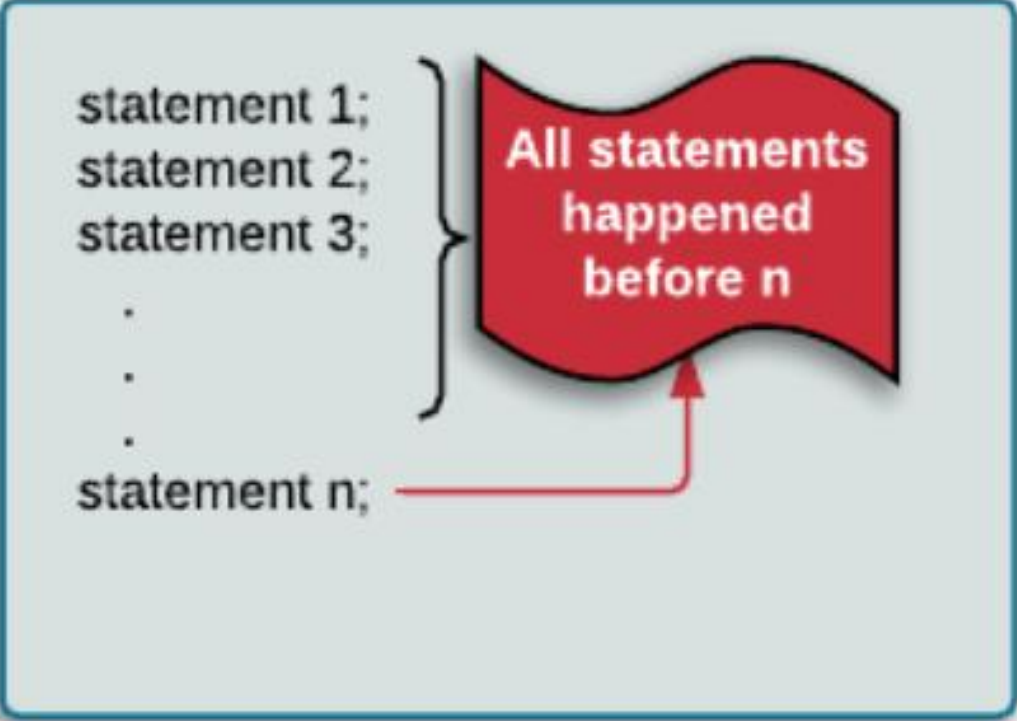




**If X should always run before Y
then X should have "happens
before" relation with Y.**

Single Thread rule

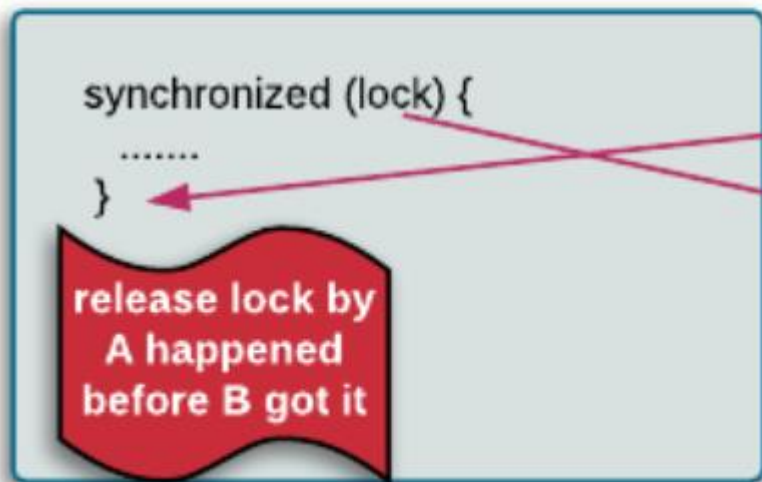
statement 1;
statement 2;
statement 3;
.
.
.
statement n;



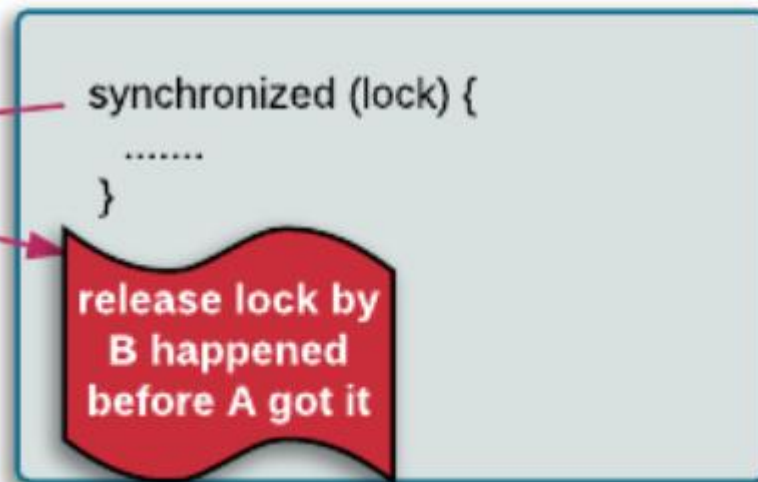
**All statements
happened
before n**

Monitor Lock rule

Thread A



Thread B



Volatile Variable Rule

Thread A

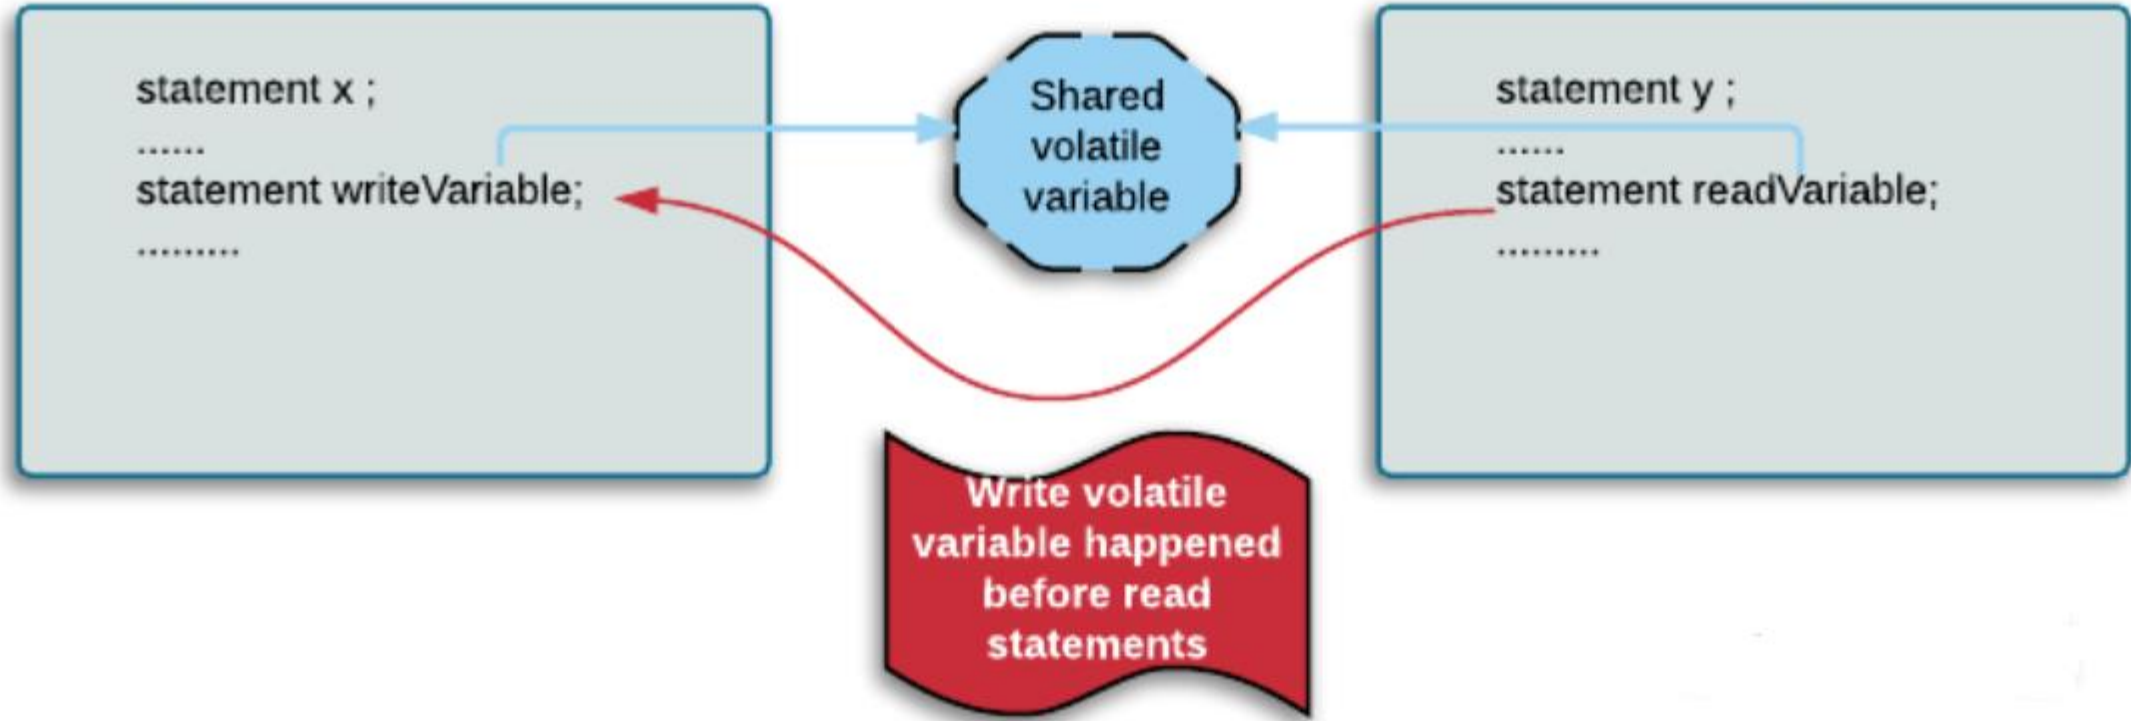
```
statement x ;  
.....  
statement writeVariable;  
.....
```

Thread B

```
statement y ;  
.....  
statement readVariable;  
.....
```

Shared
volatile
variable

Write volatile
variable happened
before read
statements



Thread Join rule

Thread A

```
.....  
Thread threadB = new Thread(..);  
.....  
threadB.start( );  
.....  
threadB.join();  
statement 1;  
.....
```

Thread B

```
.....  
public void run( ){  
    statement 1;  
    .....  
    .....  
}  
.....
```

finishing of run
method of B
happened before
statement 1


Conditions for correct use of volatile

- You can use volatile variables instead of locks only under a restricted set of circumstances. Both of the following criteria must be met for volatile variables to provide the desired thread-safety:
- Writes to the variable do not depend on its current value.
- The variable does not participate in invariants with other variables.

Listing 1. Non-thread-safe number range class

```
1  @NotThreadSafe
2  public class NumberRange {
3      private int lower, upper;
4
5      public int getLower() { return lower; }
6      public int getUpper() { return upper; }
7
8      public void setLower(int value) {
9          if (value > upper)
10             throw new IllegalArgumentException(...);
11             lower = value;
12     }
13
14     public void setUpper(int value) {
15         if (value < lower)
16             throw new IllegalArgumentException(...);
17         upper = value;
18     }
19 }
```

Performance considerations



Разликата м/у deadlock vs livelock


- Deadlock („мъртва хватка“) е ситуация, в която две или повече действия чакат друго да завърши изпълнението си, но това не се случва
- Databases -> „Мъртва хватка“ се случва когато 2 процеса искате да се подновят данни от 2 реда, но в обратна последователност. Напр. Процес А подновява ред 1 и тогава ред 2, като в същото време процес Б подновява ред 2 и след това ред 1!!!
- ОС -> „мъртва хватка“ е ситуация, която се случва, когато един процес или нишка влезе в състояние “waiting”, за да изчака необходимия му ресурс да се освободи от друг процес, но този процес от друга страна също е в състояние “waiting”, защото изчаква друг процес

livelock

- Една нишка, често действа като отговор на действието на друга нишка
- Ако друга нишка действа като отговор на действието на предходна нишка -> livelock!!!!
- Livelocked threads – това са нишки, които са не могат да изпълнят поради някоя причина. Такива нишки не са блокирани -> те са по скоро са в състояние “busy”
- Подобно на ситуацията когато двама човека се опитват да минат през коридор: А се премества на ляво за да пропусне В, докато В се премества на дясно, за да пусне А. Те се блокират един друг с тяхното действие.

Синхронизиране на нишки

Volatile



Volatile



Counter remained 1 instead of 2

~ we should make sure the threads are going to wait
for each other to finish the given task on the variables !!!

Как работи *multithreading*?

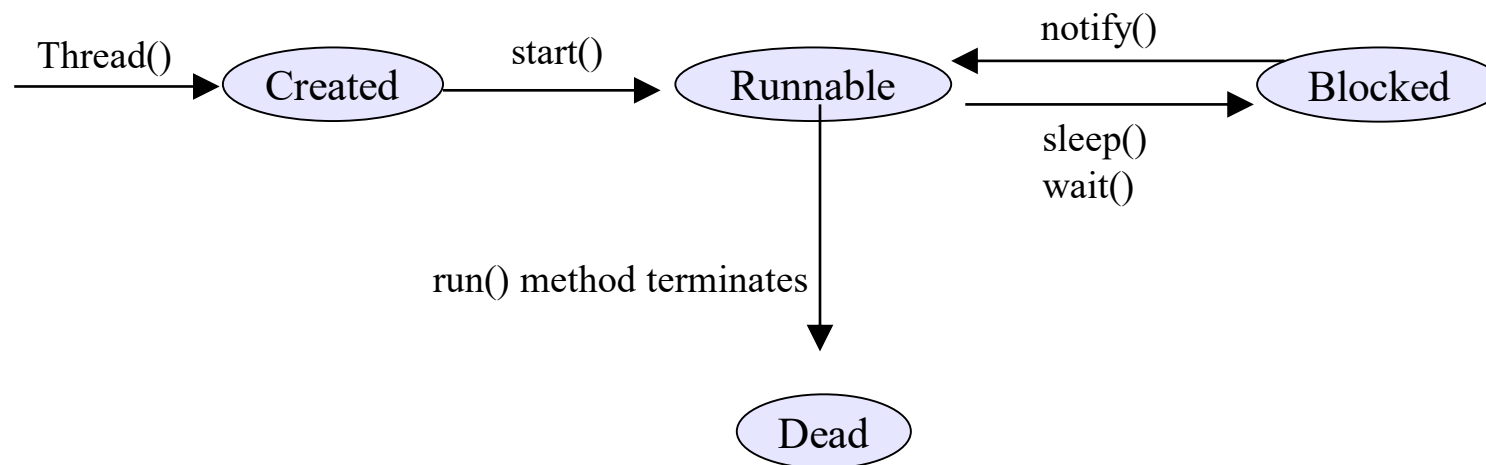
- Всяка нишка има свое собствен "context"
 - Всеки context съдържа виртуални регистри и отделен стек
- Така наречения "scheduler" взима решение коя нишка да се изпълнява в даден момент.
 - За виртуалната памет може да има отделен scheduler
 - Много ОСи директно поддържат multithreading, виртуалната памет може да използва системния диспечър(scheduler) за поддръжка на нишките
- Диспечърът(scheduler) поддържа списък с готови нишки (които са записани в run queue), както и списък от нишки в състояние waiting (wait queue)
- Всяка нишка има приоритет. Диспечърът избира за изпълнение тези нишки с най-висок от run queue
 - Забележете: програмистът не може да знае колко са нишките, готови за изпълнение и обикновено, броят на нишките, както и тяхното изпълнение ще различно върху различни платформи.

Реализация на нишки в Java

- Някои езици за програмиране директно поддържат нишки
 - Поддържат add-on thread режим
 - Режимът add on thread в повечето случаи е сложен за използване
- В JVM има отделни нишки от ОС
 - В JVM се използва garbage collection
- Нишките се представят от класа Thread.
 - Обект от класа има в себе си състояние на нишка
 - Осигурява методи като interrupt, start, sleep, yield, wait
- Методът main се изпълнява от главната нишка.
 - Когато приложението изисква повече нишки, то те трябва допълнително да се създават.


Състояния на нишките

- Нишките могат да бъдат в едно от 4-те състояния
 - Created, Running, Blocked, и Dead
- Състоянието на нишките се променя в зависимост от:
 - Методите за управление на нишките :start, sleep, yield, wait, notify
 - Спиране на изпълнението на програмата



Как се стартира Thread?

- Извиква се метода **run()**
 - run() се изпълнява, когато метода start() се извиква от нишката
- Нишката ще се терминира ако метода run() е терминиран
 - За да не се терминира нишката, метода run() трябва да не се прекрати
 - Методът run() често има в себе си безкраен цикъл, за да не се спре нишката
- Една нишка извиква друга, когато тя извика метода start на другата нишка.
 - Последователността от събития може да обърка тези, които се по-запознати с едно-нишковото програмиране.




Creating your own Threads

- The obvious way to create your own threads is to subclass the Thread class and then override the run() method
 - This is the easiest way to do it
 - It is not the recommended way to do it.
- Because threads are usually associated with a task, the object which provides the run method is usually a subclass of some other class
 - If it inherits from another class, it cannot inherit from Thread.
- The solution is provided by an interface called Runnable.
 - Runnable defines one method - public void run()
- One of the Thread classes constructor takes a reference to a Runnable object
 - When the thread is started, it invokes the run method in the runnable object instead of its own run method.

Using Runnable

- In the example below, when the Thread object is instantiated, it is passed a reference to a "Runnable" object
 - The Runnable object must implement a method called "run"
- When the thread object receives a start message, it checks to see if it has a reference to a Runnable object:
 - If it does, it runs the "run" method of that object
 - If not, it runs its own "run" method




Example Code

```
public class Test implements Runnable
{
    private Thread theThread;

    public void start()
    {
        if (theThread == null)
        {
            theThread = new Thread(this);
            theThread.start();
        }
    }

    public void run()
    {
        // This method runs in its
        // own thread
    }
}
```



Properly Terminating Threads

- In Java 1.1, the Thread class had a stop() method
 - One thread could terminate another by invoking its stop() method.
 - However, using stop() could lead to deadlocks
 - The stop() method is now deprecated. DO NOT use the stop method to terminate a thread
- The correct way to stop a thread is to have the run method terminate
 - Add a boolean variable which indicates whether the thread should continue or not
 - Provide a set method for that variable which can be invoked by another thread

Terminating Thread Example

```
public class Test implements Runnable
{
    private Thread theThread;
    private boolean stopThread = false;

    public void start()
    {
        if (theThread == null)
        {
            theThread = new Thread(this);
            theThread.start();
        }
    }

    public void setStopThread(boolean aValue)
    {
        stopThread = aValue;
    }

    public void run()
    {
        while(true)
        {
            if (stopThread)
                break;

            // ....
        }
    }
}
```

Creating Multiple Threads

- The previous example illustrates a Runnable class which creates its own thread when the start method is invoked.
- If one wished to create multiple threads, one could simple create multiple instances of the Runnable class and send each object a start message
 - Each instance would create its own thread object
- Is the a maximum number of threads which can be created?
 - There is no defined maximum in Java.
 - If the VM is delegating threads to the OS, then this is platform dependent.
 - A good rule of thumb for maximum thread count is to allow 2Mb of ram for each thread
 - Although threads share the same memory space, this can be a reasonable estimate of how many threads your machine can handle.

Thread Priorities

- Every thread is assigned a priority (between 1 and 10)
 - The default is 5
 - The higher the number, the higher the priority
 - Can be set with `setPriority(int aPriority)`
- The standard mode of operation is that the scheduler executes threads with higher priorities first.
 - This simple scheduling algorithm can cause problems. Specifically, one high priority thread can become a "CPU hog".
 - A thread using vast amounts of CPU can share CPU time with other threads by invoking the `yield()` method on itself.
- Most OSes do not employ a scheduling algorithm as simple as this one
 - Most modern OSes have thread aging
 - The more CPU a thread receives, the lower its priority becomes
 - The more a thread waits for the CPU, the higher its priority becomes
 - Because of thread aging, the effect of setting a thread's priority is dependent on the platform

Yield() and Sleep()

- Sometimes a thread can determine that it has nothing to do
 - Sometimes the system can determine this. ie. waiting for I/O
- When a thread has nothing to do, it should not use CPU
 - This is called a busy-wait.
 - Threads in busy-wait are busy using up the CPU doing nothing.
 - Often, threads in busy-wait are continually checking a flag to see if there is anything to do.
- It is worthwhile to run a CPU monitor program on your desktop
 - You can see that a thread is in busy-wait when the CPU monitor goes up (usually to 100%), but the application doesn't seem to be doing anything.
- Threads in busy-wait should be moved from the Run queue to the Wait queue so that they do not hog the CPU
 - Use yield() or sleep(time)
 - Yield simply tells the scheduler to schedule another thread
 - Sleep guarantees that this thread will remain in the wait queue for the specified number of milliseconds.

Concurrent Access to Data

- Those familiar with databases will understand that concurrent access to data can lead to data integrity problems
 - Specifically, if two sources attempt to update the same data at the same time, the result of the data can be undefined.
 - The outcome is determined by how the scheduler schedules the two sources.
 - Since the schedulers activities cannot be predicted, the outcome cannot be predicted
- Databases deal with this mechanism through "locking"
 - If a source is going to update a table or record, it can lock the table or record until such time that the data has been successfully updated.
 - While locked, all access is blocked except to the source which holds the lock.
- Java has the equivalent mechanism. It is called synchronization
 - Java has a keyword called synchronized

Synchronization

- In Java, every object has a lock
 - To obtain the lock, you must synchronize with the object
- The simplest way to use synchronization is by declaring one or more methods to be synchronized
 - When a synchronized method is invoked, the calling thread attempts to obtain the lock on the object.
 - if it cannot obtain the lock, the thread goes to sleep until the lock becomes available
 - Once the lock is obtained, no other thread can obtain the lock until it is released. ie, the synchronized method terminates
 - When a thread is within a synchronized method, it knows that no other synchronized method can be invoked by any other thread
 - Therefore, it is within synchronized methods that critical data is updated

Providing Thread Safe Access to Data

- If an object contains data which may be updated from multiple thread sources, the object should be implemented in a thread-safe manner
 - All access to critical data should only be provided through synchronized methods (or synchronized blocks).
 - In this way, we are guaranteed that the data will be updated by only one thread at a time.

```
public class SavingsAccount
{
    private float balance;

    public synchronized void withdraw(float anAmount)
    {
        if ((anAmount>0.0) && (anAmount<=balance))
            balance = balance - anAmount;
    }

    public synchronized void deposit(float anAmount)
    {
        if (anAmount>0.0)
            balance = balance + anAmount;
    }
}
```

Thread Safety Performance Issues

- However, there is an overhead associated with synchronization
 - Many threads may be waiting to gain access to one of the object's synchronized methods
 - The object remains locked as long as a thread is within a synchronized method.
 - Ideally, the method should be kept as short as possible.
- Another solution is to provide synchronization on a block of code instead of the entire method
 - In this case, the object's lock is only held for the time that the thread is within the block.
 - The intent is that we only lock the region of code which requires access to the critical data. Any other code within the method can occur without the lock.
 - In high load situations where multiple threads are attempting to access critical data, this is by far a much better implementation.

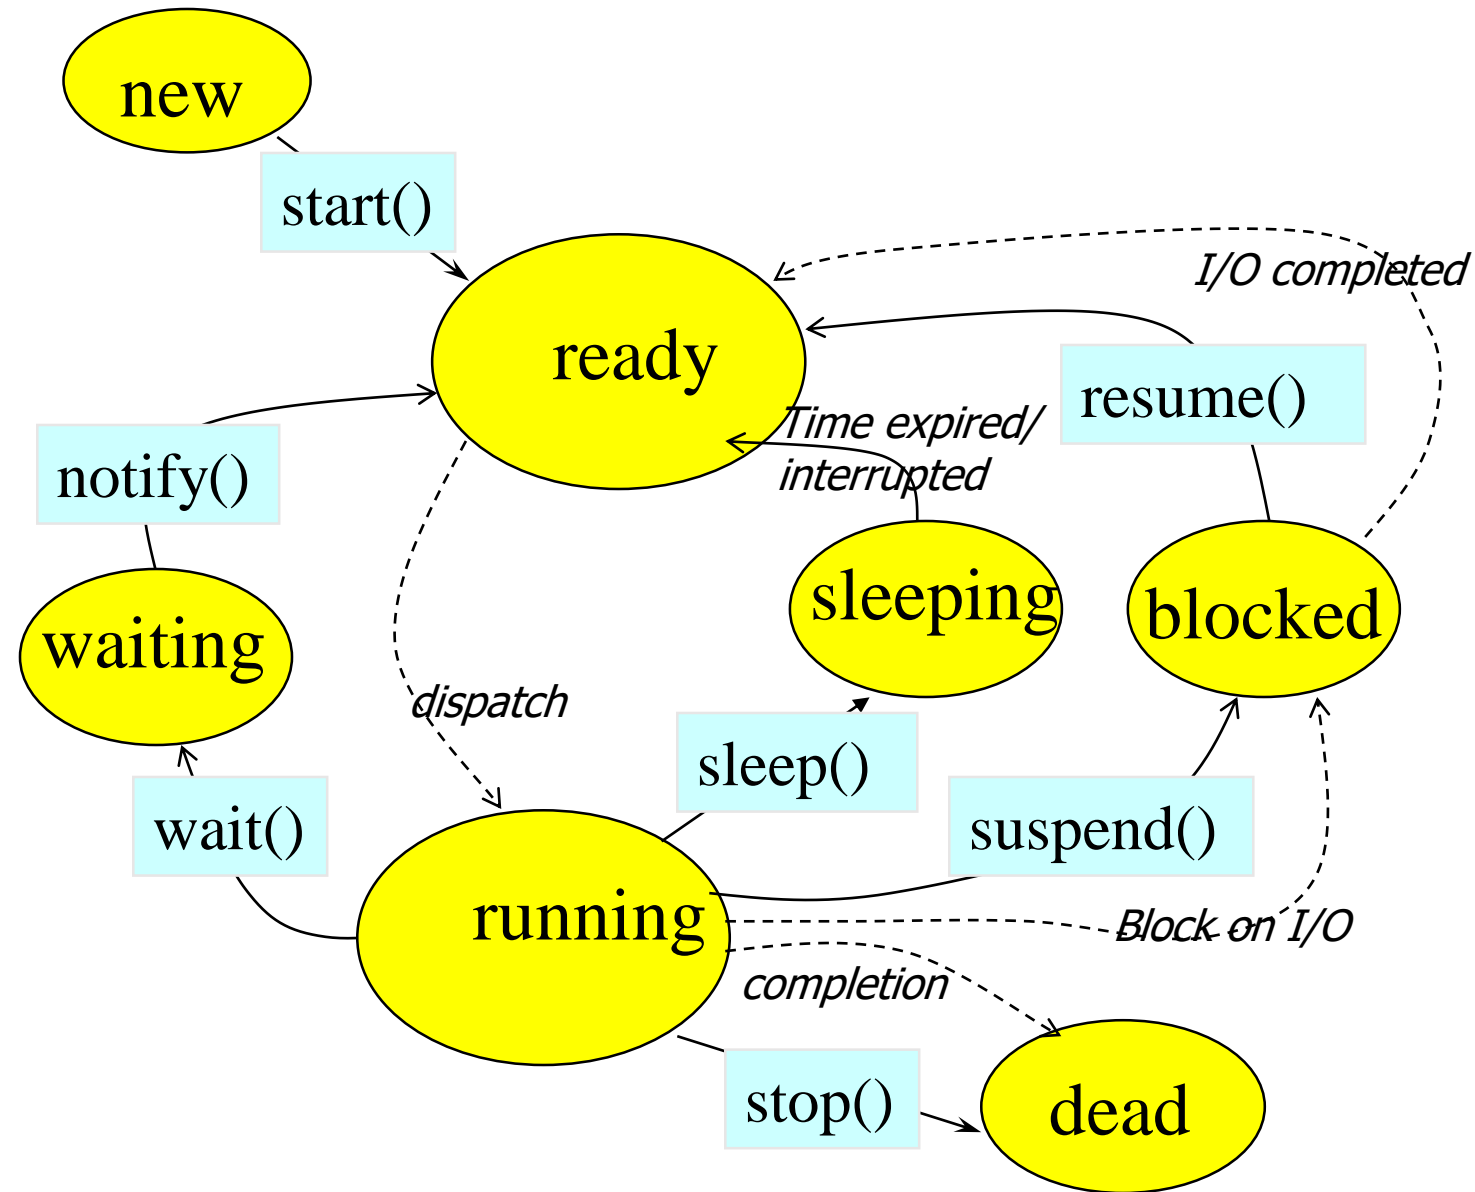
Block Synchronization

```
public class SavingsAccount
{
    private float balance;

    public void withdraw(float anAmount)
    {
        if (anAmount<0.0)
            throw new IllegalArgumentException("Withdraw amount negative");
        synchronized(this)
        {
            if (anAmount<=balance)
                balance = balance - anAmount;
        }
    }

    public void deposit(float anAmount)
    {
        if (anAmount<0.0)
            throw new IllegalArgumentException("Deposit amount negative");
        synchronized(this)
        {
            balance = balance + anAmount;
        }
    }
}
```


Life Cycle of Thread



A Program with Three Java Threads

- Write a program that creates 3 threads

Three threads example

```
class A extends Thread
{
    public void run()
    {
        for(int i=1;i<=5;i++)
        {
            System.out.println("\t From ThreadA: i= "+i);
        }
        System.out.println("Exit from A");
    }
}
```

```
class B extends Thread
{
    public void run()
    {
        for(int j=1;j<=5;j++)
        {
            System.out.println("\t From ThreadB: j= "+j);
        }
        System.out.println("Exit from B");
    }
}
```

```
class C extends Thread
```

```
{  
    public void run()  
    {  
        for(int k=1;k<=5;k++)  
        {  
            System.out.println("\t From ThreadC: k= "+k);  
        }  
  
        System.out.println("Exit from C");  
    }  
}
```

```
class ThreadTest
```

```
{  
    public static void main(String args[])  
    {  
        new A().start();  
        new B().start();  
        new C().start();  
    }  
}
```

Run 1

- [raj@mundroo] threads [1:76] java ThreadTest

From ThreadA: i= 1

From ThreadA: i= 2

From ThreadA: i= 3

From ThreadA: i= 4

From ThreadA: i= 5

Exit from A

From ThreadC: k= 1

From ThreadC: k= 2

From ThreadC: k= 3

From ThreadC: k= 4

From ThreadC: k= 5

Exit from C

From ThreadB: j= 1

From ThreadB: j= 2

From ThreadB: j= 3

From ThreadB: j= 4

From ThreadB: j= 5

Exit from B

Run2

- [raj@mundroo] threads [1:77] java ThreadTest

From ThreadA: i= 1

From ThreadA: i= 2

From ThreadA: i= 3

From ThreadA: i= 4

From ThreadA: i= 5

From ThreadC: k= 1

From ThreadC: k= 2

From ThreadC: k= 3

From ThreadC: k= 4

From ThreadC: k= 5

Exit from C

From ThreadB: j= 1

From ThreadB: j= 2

From ThreadB: j= 3

From ThreadB: j= 4

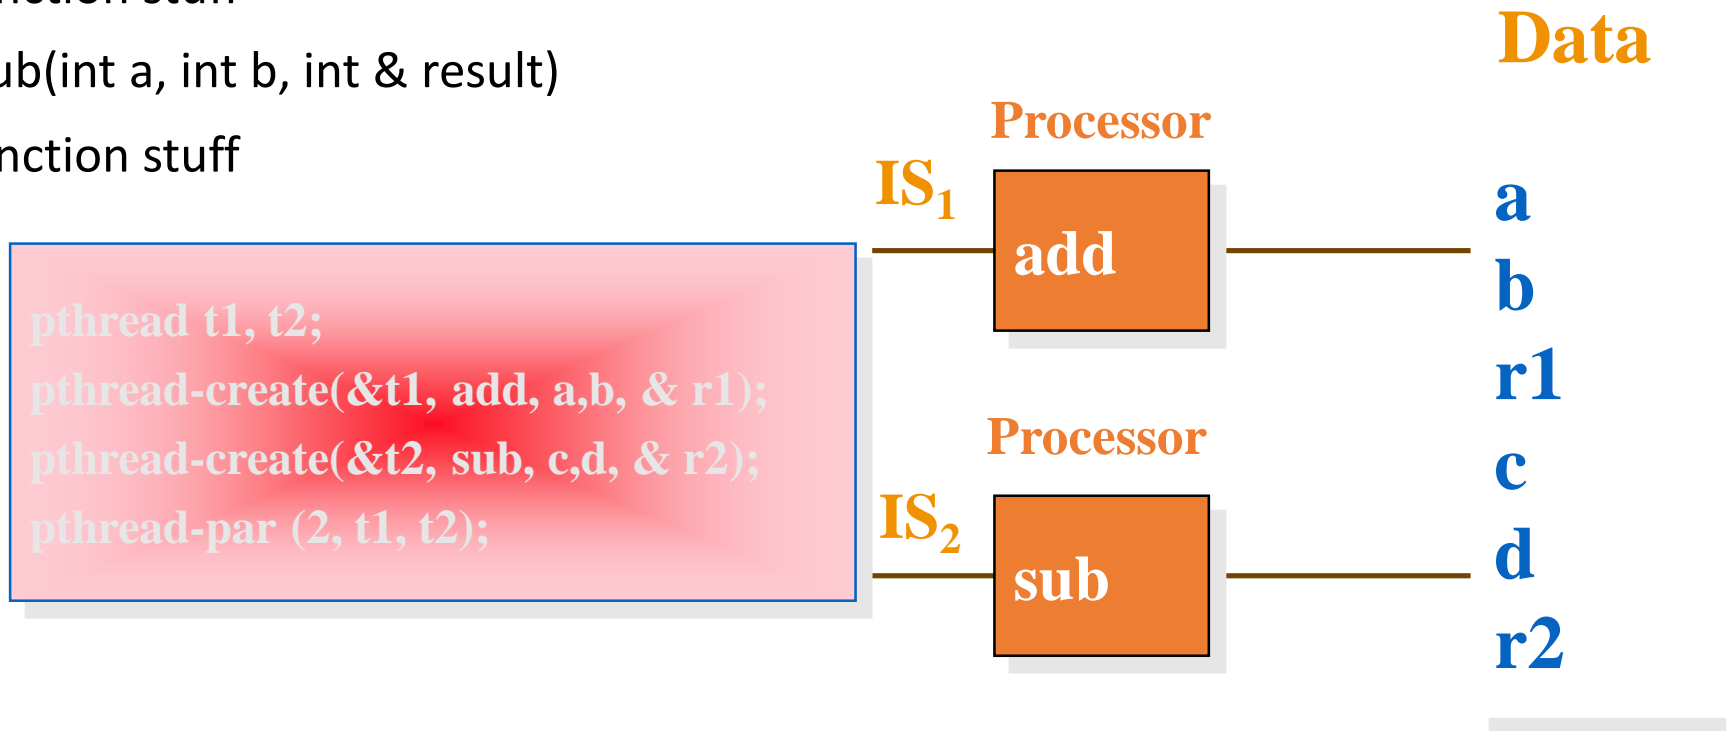
From ThreadB: j= 5

Exit from B

Exit from A

Process Parallelism

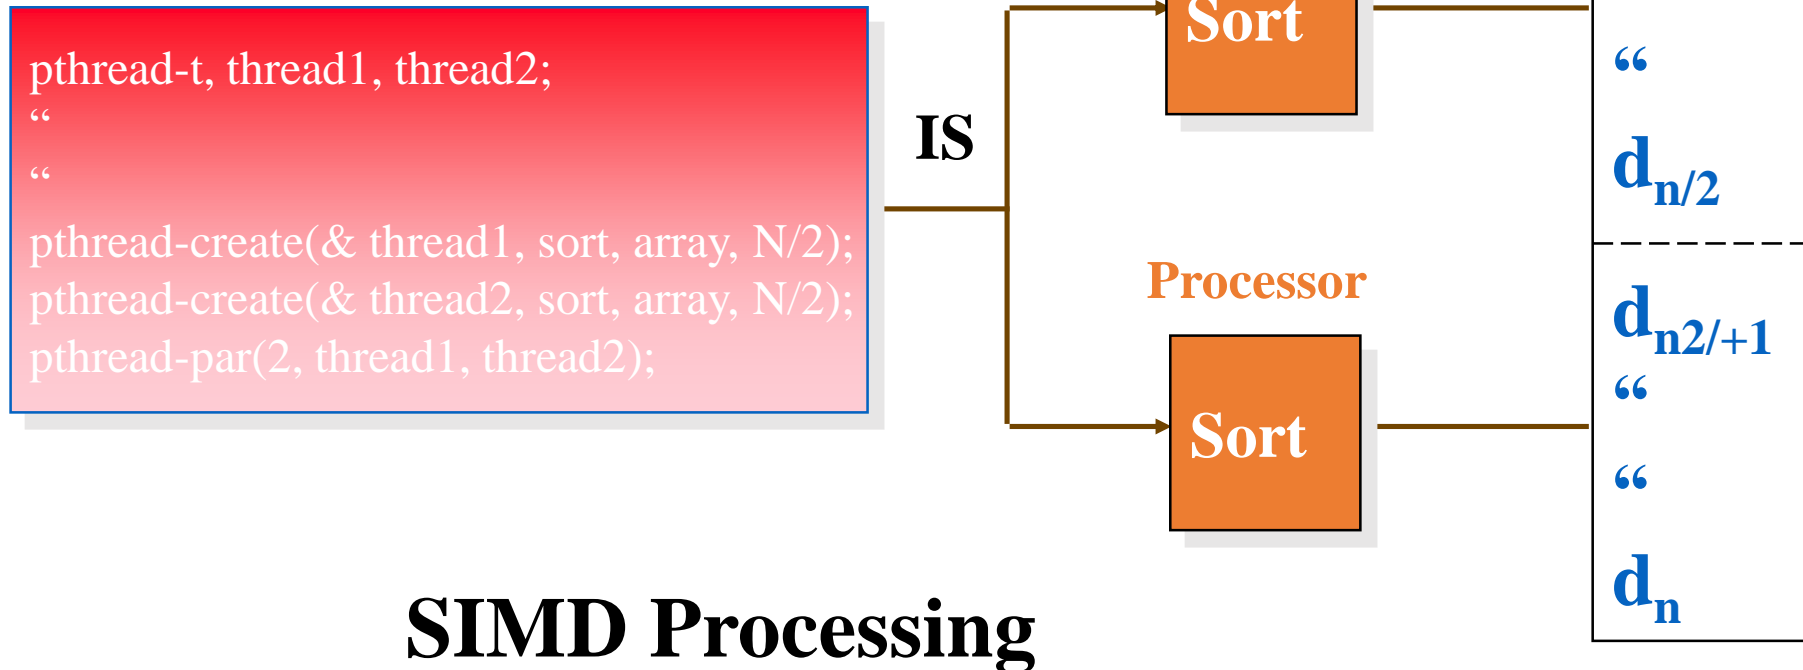
- `int add (int a, int b, int & result)`
- `// function stuff`
- `int sub(int a, int b, int & result)`
- `// function stuff`



MISD and MIMD Processing

Data Parallelism

- `sort(int *array, int count)`
- `//.....`
- `//.....`



Thread Priority

- In Java, each thread is assigned priority, which affects the order in which it is scheduled for running. The threads so far had same default priority (NORM_PRIORITY) and they are served using FCFS policy.
 - Java allows users to change priority:
 - ThreadName.setPriority(intNumber)
 - MIN_PRIORITY = 1
 - NORM_PRIORITY=5
 - MAX_PRIORITY=10

Thread Priority Example

```
class A extends Thread
{
    public void run()
    {
        System.out.println("Thread A started");
        for(int i=1;i<=4;i++)
        {
            System.out.println("\t From ThreadA: i= "+i);
        }
        System.out.println("Exit from A");
    }
}

class B extends Thread
{
    public void run()
    {
        System.out.println("Thread B started");
        for(int j=1;j<=4;j++)
        {
            System.out.println("\t From ThreadB: j= "+j);
        }
        System.out.println("Exit from B");
    }
}
```

Thread Priority Example

```
class C extends Thread
{
    public void run()
    {
        System.out.println("Thread C started");
        for(int k=1;k<=4;k++)
        {
            System.out.println("\t From ThreadC: k= "+k);
        }
        System.out.println("Exit from C");
    }
}
```

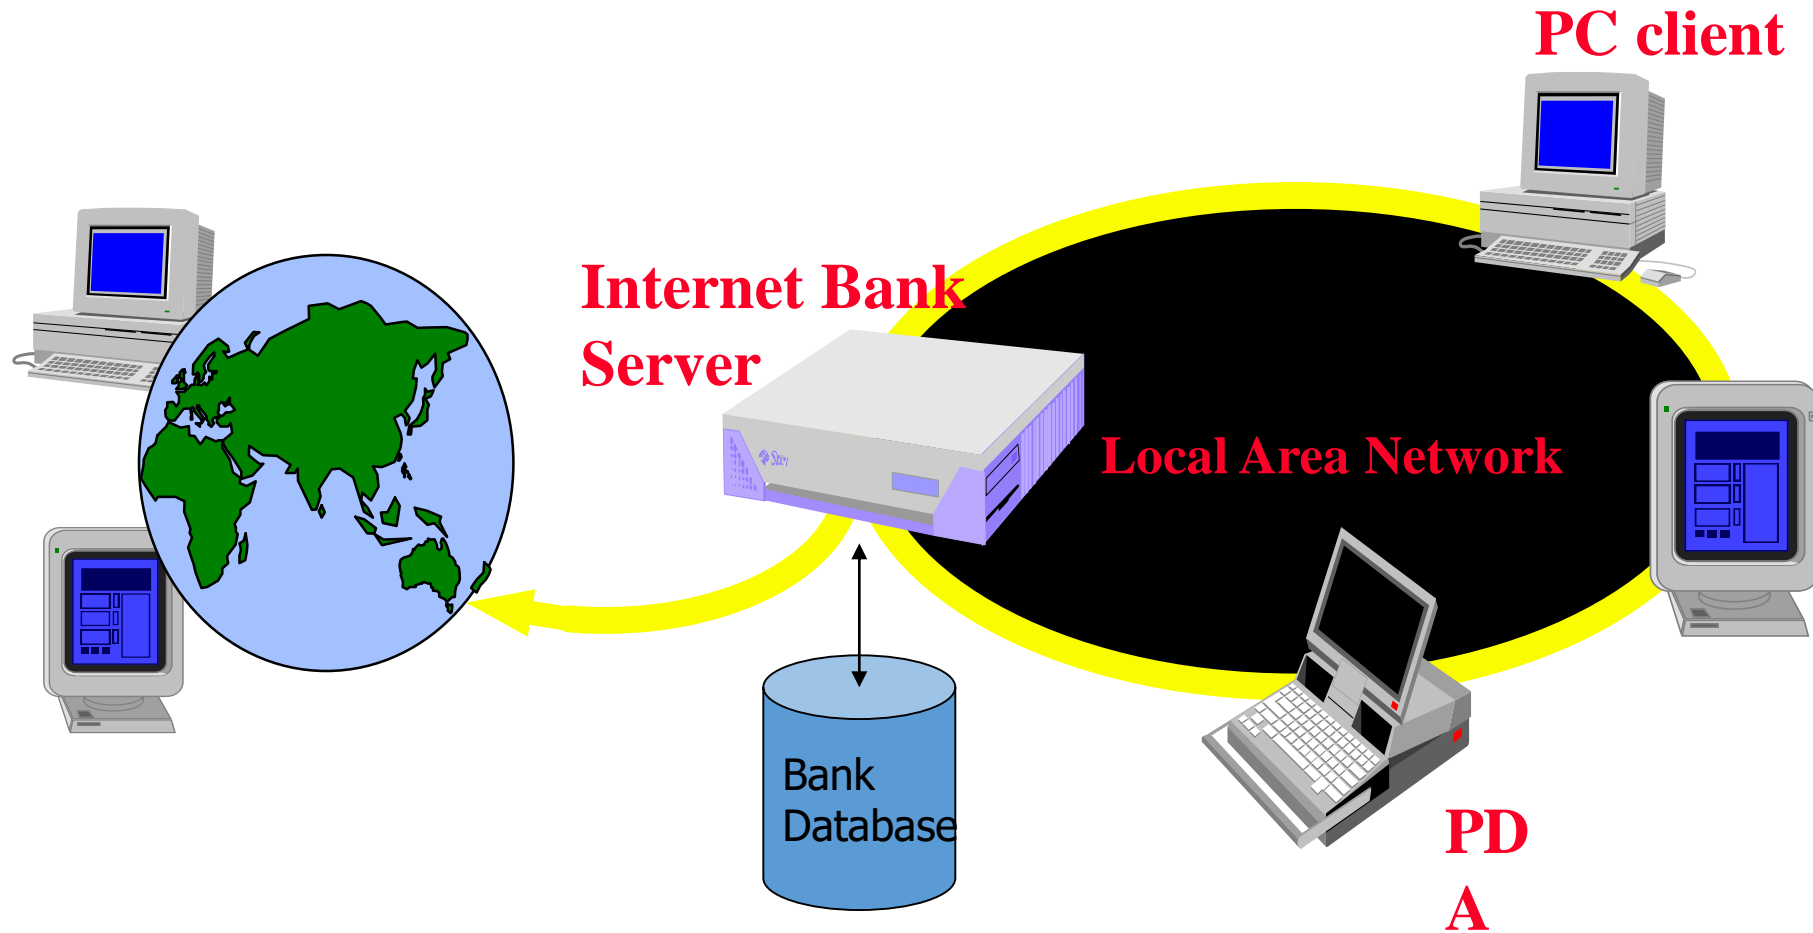
77

```
class ThreadPriority
{
    public static void main(String args[])
    {
        A threadA=new A();
        B threadB=new B();
        C threadC=new C();
        threadC.setPriority(Thread.MAX_PRIORITY);
        threadB.setPriority(threadA.getPriority()+1);
        threadA.setPriority(Thread.MIN_PRIORITY);
        System.out.println("Started Thread A");
        threadA.start();
        System.out.println("Started Thread B");
        threadB.start();
        System.out.println("Started Thread C");
        threadC.start();
        System.out.println("End of main thread");
    }
}
```

Accessing Shared Resources

- Applications Access to Shared Resources need to be coordinated.
 - Printer (two person jobs cannot be printed at the same time)
 - Simultaneous operations on your bank account.
 - Can the following operations be done at the same time on the same account?
 - Deposit()
 - Withdraw()
 - Enquire()

Online Bank: Serving Many Customers and Operations



Shared Resources



- If one thread tries to read the data and other thread tries to update the same data, it leads to inconsistent state.
- This can be prevented by synchronising access to the data.
- Use “Synchronized” method:
 - public **synchronized** void update()
 - {
 - ...
 - }

the driver: 3rd Threads sharing the same object


```
class InternetBankingSystem {  
    public static void main(String [] args ) {  
        Account accountObject = new Account ();  
        Thread t1 = new Thread(new MyThread(accountObject));  
        Thread t2 = new Thread(new YourThread(accountObject));  
        Thread t3 = new Thread(new HerThread(accountObject));  
        t1.start();  
        t2.start();  
        t3.start();  
        // DO some other operation  
    } // end main()  
}
```

Shared account object between 3 threads

```
class MyThread implements Runnable {  
    Account account;  
    public MyThread (Account s) { account = s;}  
    public void run() { account.deposit(); }  
} // end class MyThread
```

```
class YourThread implements Runnable {  
    Account account;  
    public YourThread (Account s) { account = s;}  
    public void run() { account.withdraw(); }  
} // end class YourThread
```

```
class HerThread implements Runnable {  
    Account account;  
    public HerThread (Account s) { account = s; }  
    public void run() { account.enquire(); }  
} // end class HerThread
```



Monitor (shared object access): serializes operation on shared object

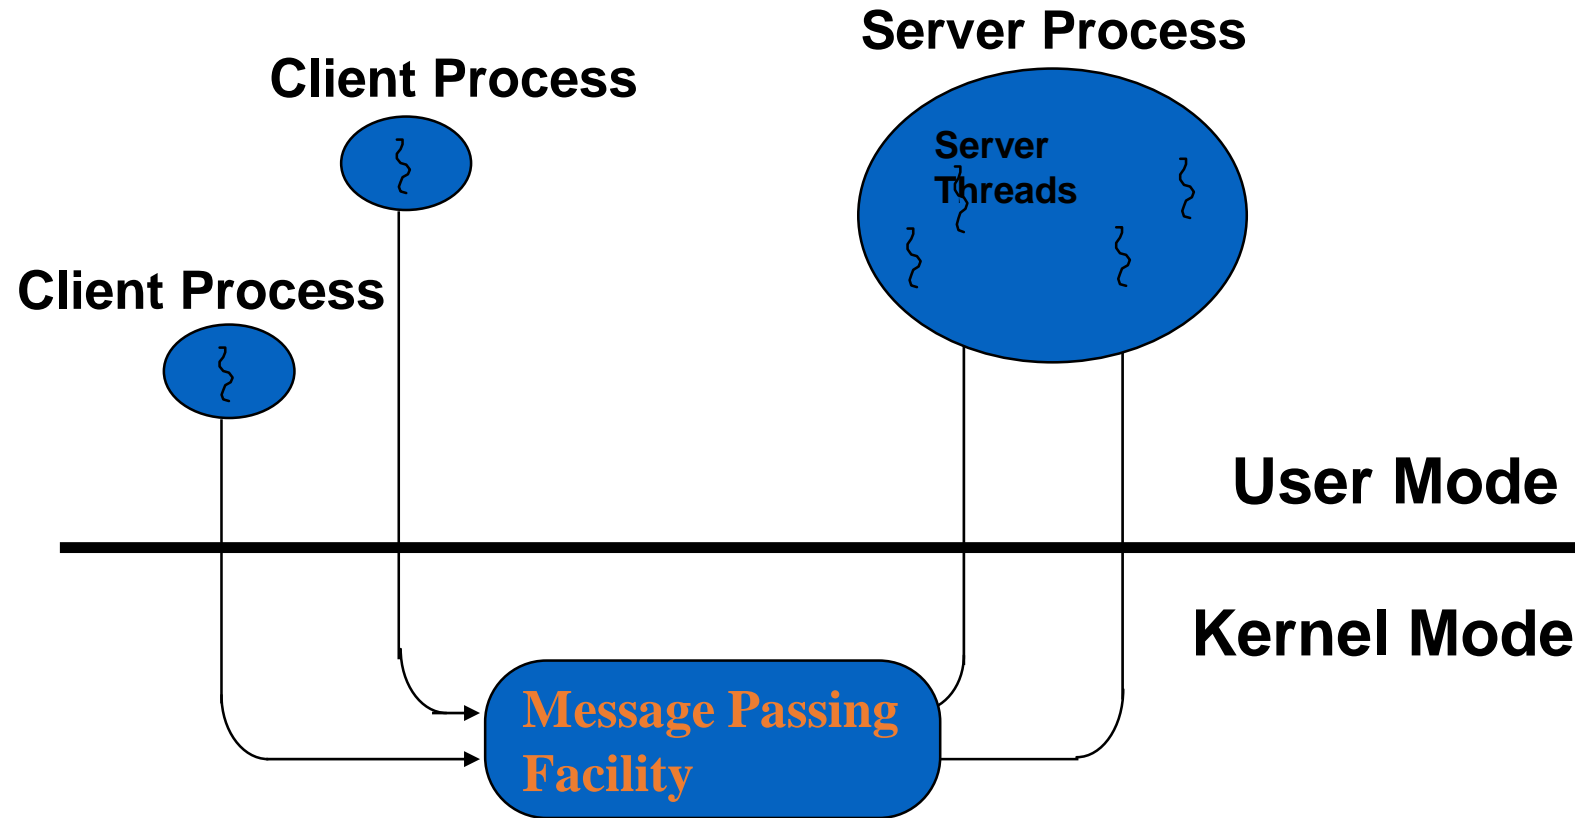
```
class Account { // the 'monitor'
    int balance;

    // if 'synchronized' is removed, the outcome is unpredictable
    public synchronized void deposit( ) {
        // METHOD BODY : balance += deposit_amount;
    }

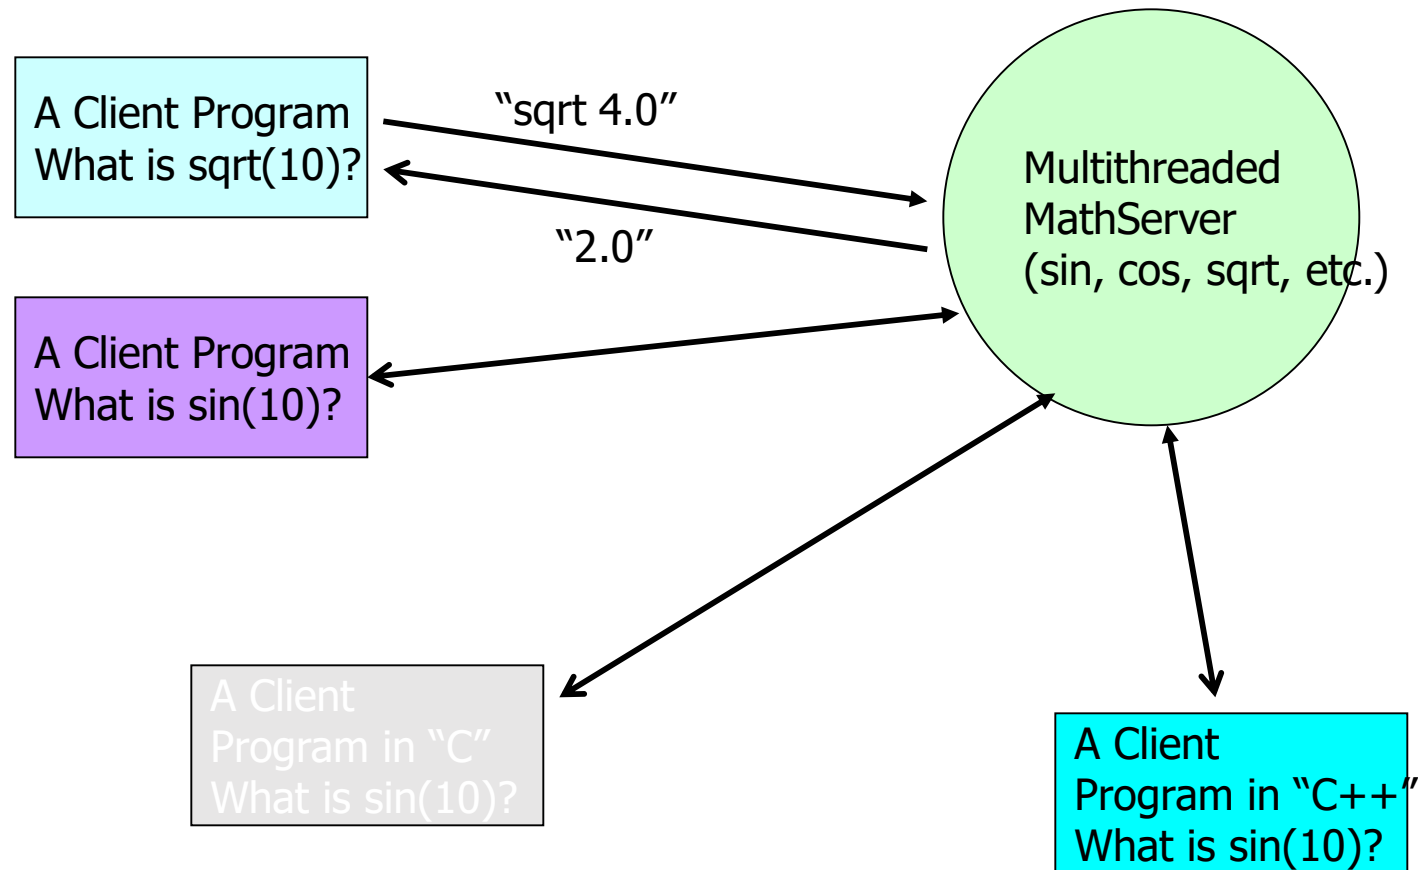
    public synchronized void withdraw( ) {
        // METHOD BODY: balance -= deposit_amount;
    }
    public synchronized void enquire( ) {
        // METHOD BODY: display balance.
    }
}
```

Multithreaded Server


Multithreaded Server



Assignment 1: Multithreaded MathServer – Demonstrates the use of Sockets and Threads



A Multithreaded Program

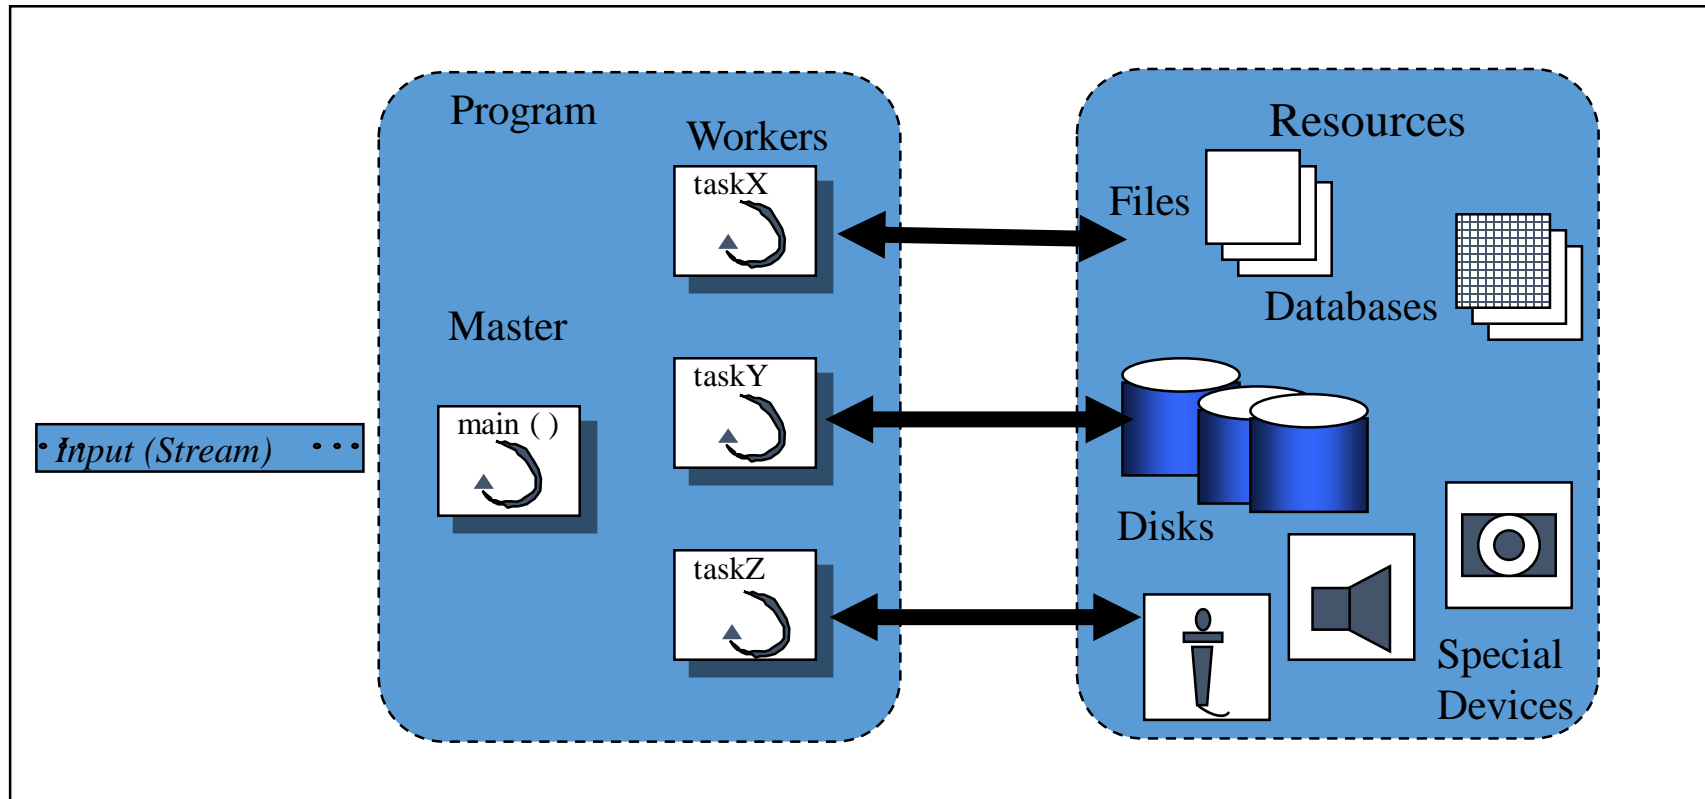


Thread Programming models


Thread concurrency/operation models

- The master/worker model
- The peer model
- A thread pipeline

The master/worker model

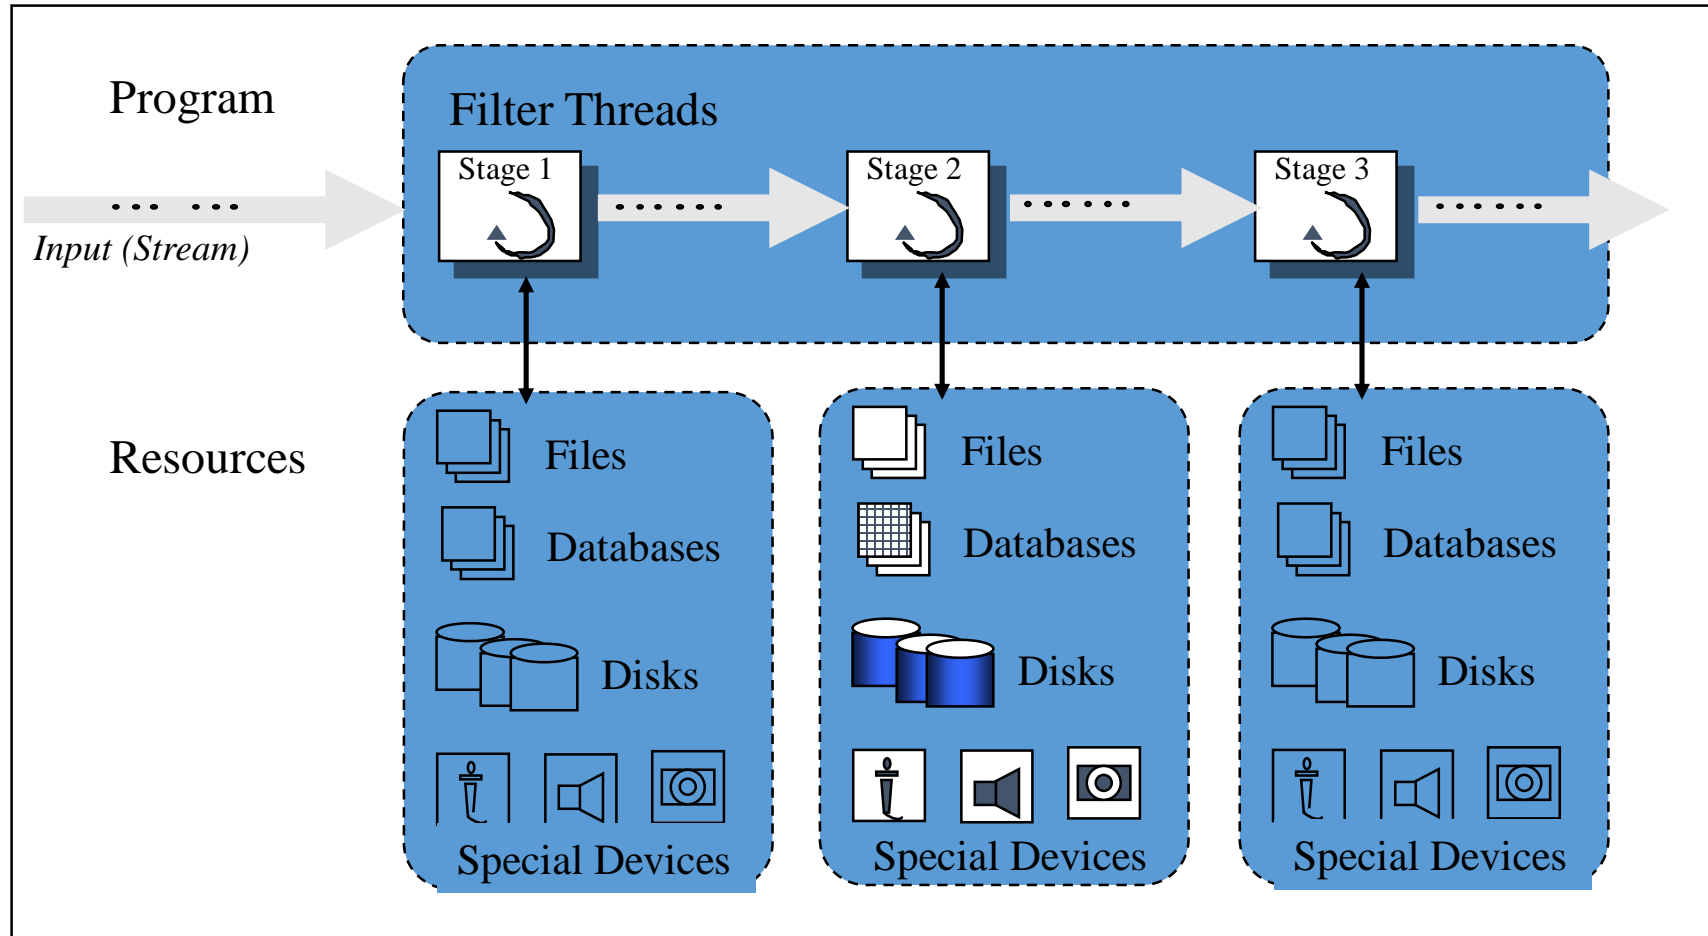


The peer model



A thread pipeline

A thread pipeline

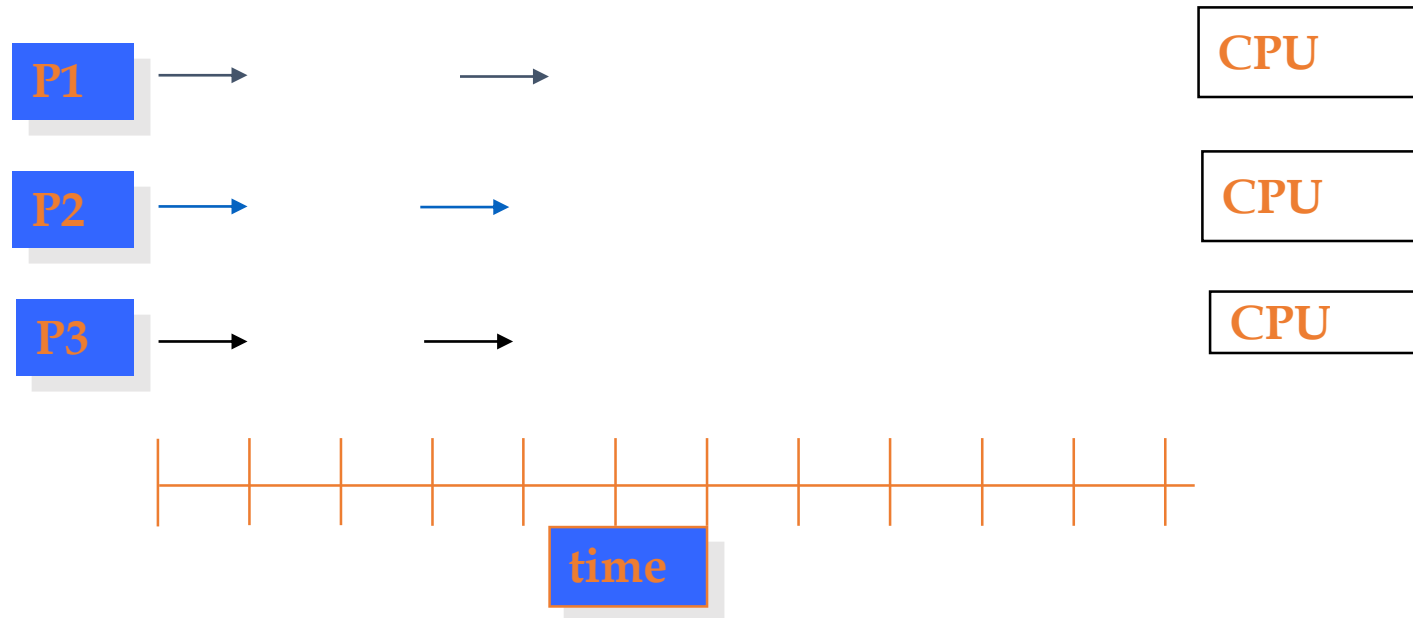


Multithreading and Multiprocessing Deployment issues

On Shared and distributed memory systems

Multithreading - Multiprocessors

Process Parallelism

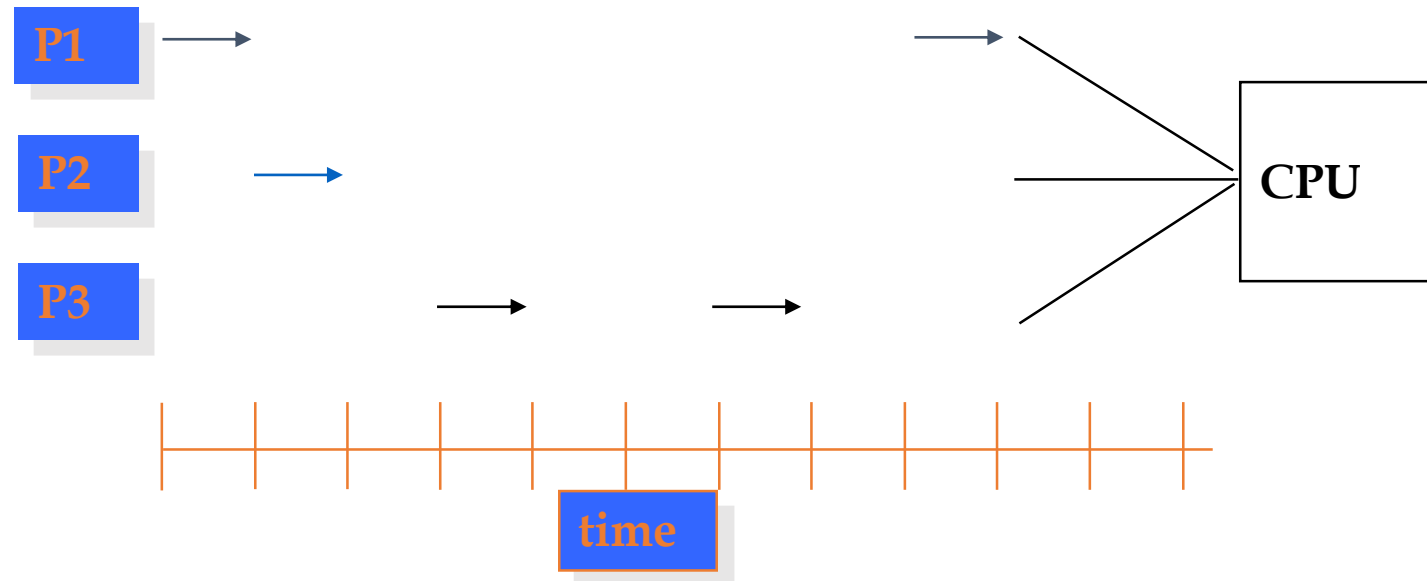


No of execution processes \leq the number of CPUs

Multithreading on Uni-processor

- Concurrency Vs Parallelism


☹ Process Concurrency



Number of Simultaneous execution units > number of CPUs

Multi-Processing (clusters & grids) and Multi-Threaded Computing

Threaded Libraries, Multi-threaded I/O



*Better Response Times in
Multiple Application
Environments*

*Higher Throughput for
Parallelizeable Applications*